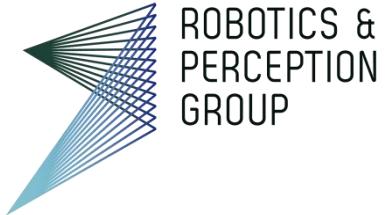




University of  
Zurich<sup>UZH</sup>



# Vision Algorithms for Mobile Robotics

## Lecture 12c Deep Learning Tutorial

Daniel Gehrig and Elia Kaufmann

<http://rpg.ifi.uzh.ch>

# Outline

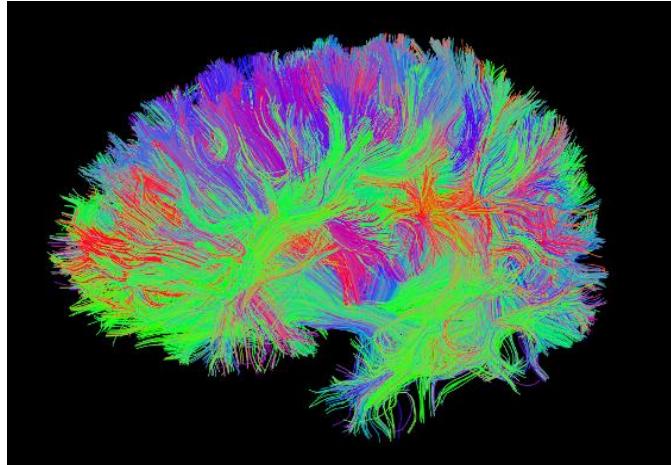
- Introduction
  - Supervised Learning
  - Unsupervised Learning
  - Applications to Computer Vision
  - Conclusions
  - Machine Learning for Drones
- 
- Relevant for the exam**

# Outline

- Introduction
- Supervised Learning
- Unsupervised Learning
- Applications to Computer Vision
- Conclusions
- Machine Learning for Drones

# The Deep Learning Revolution

Medicine



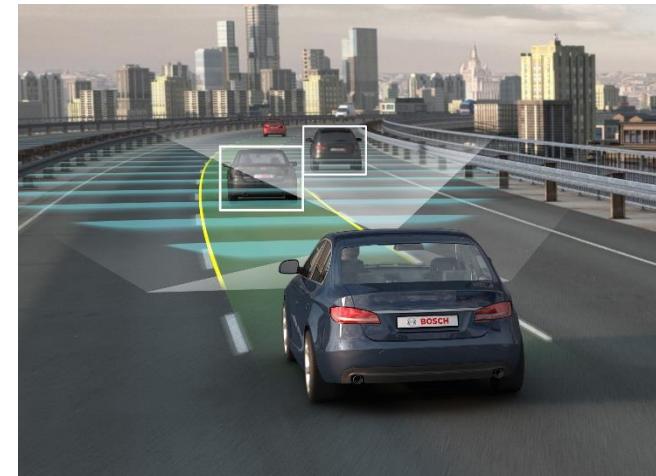
Media & Entertainment



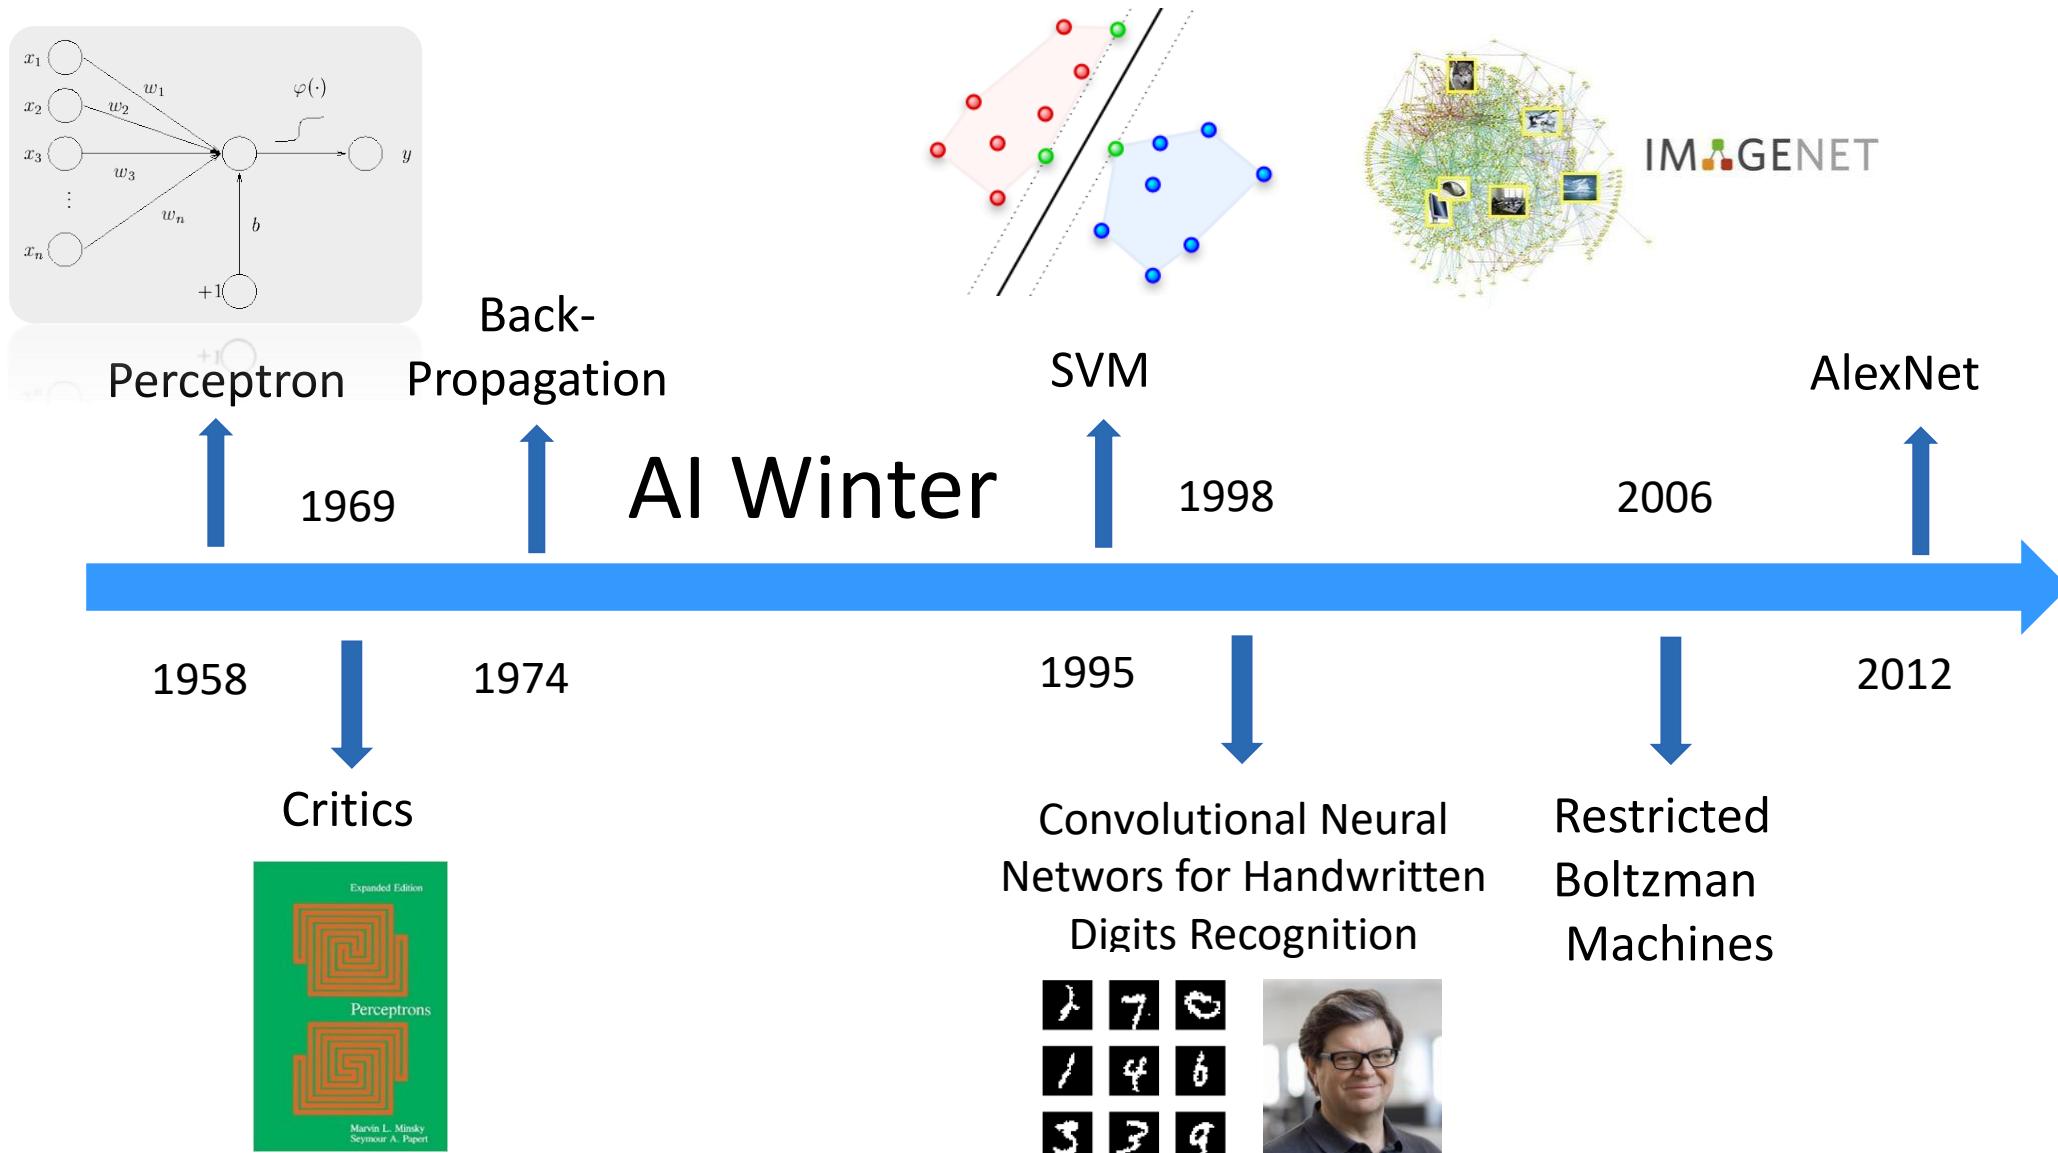
Surveillance & Security



Autonomous Driving



# Some History



# What changed?

1. Hardware Improvements



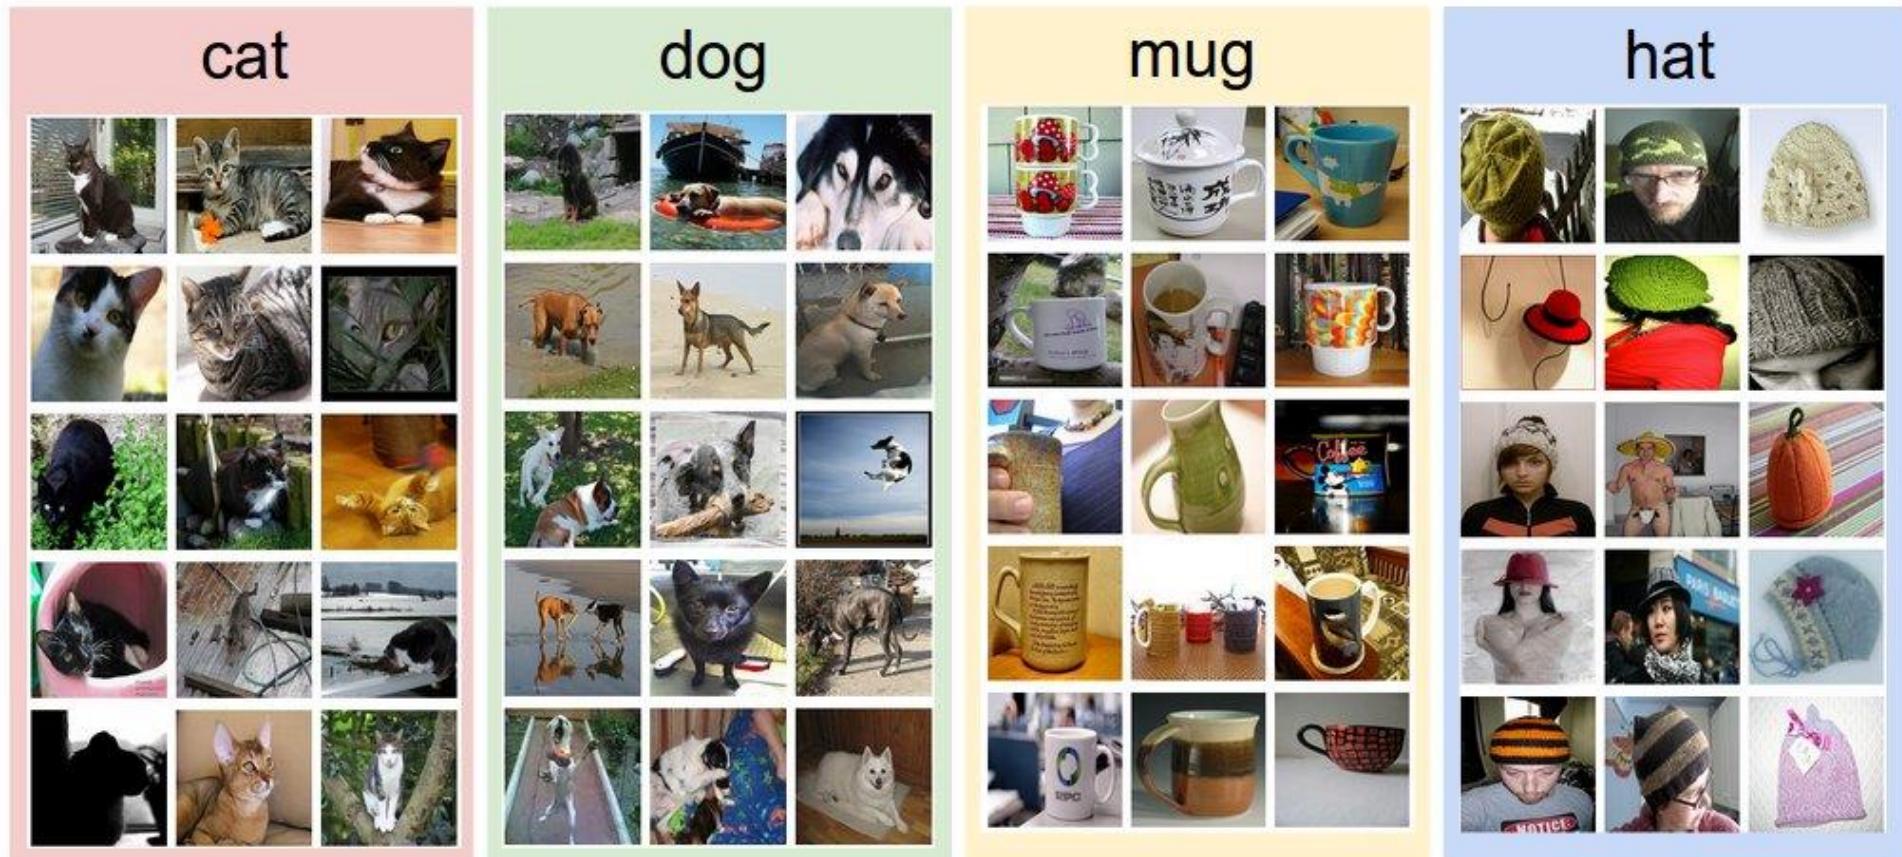
2. Big Data Available



3. Algorithmic Progress

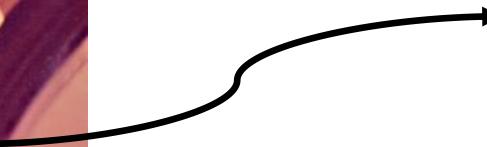
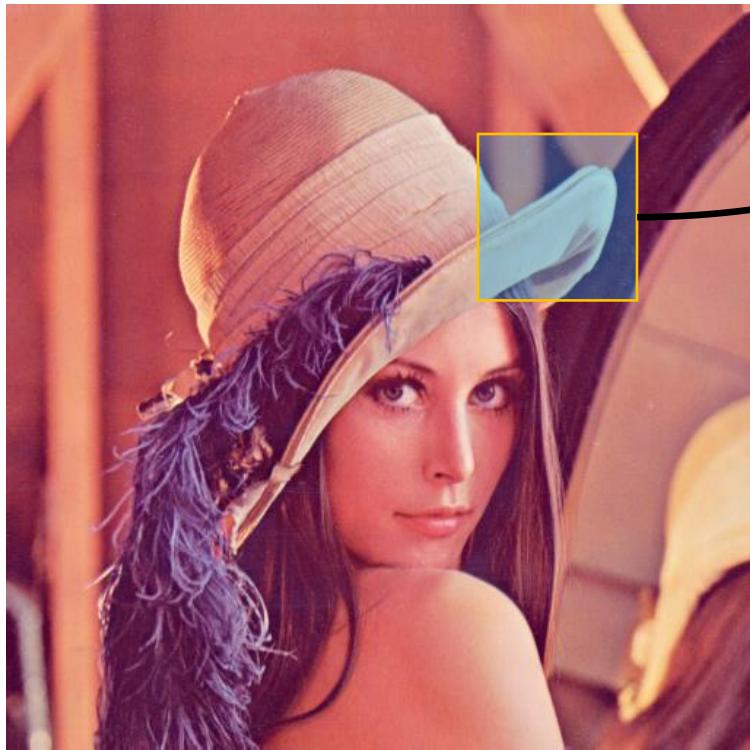
# Image Classification

Task of assigning an input image a label from a fixed set of categories.



# The semantic gap

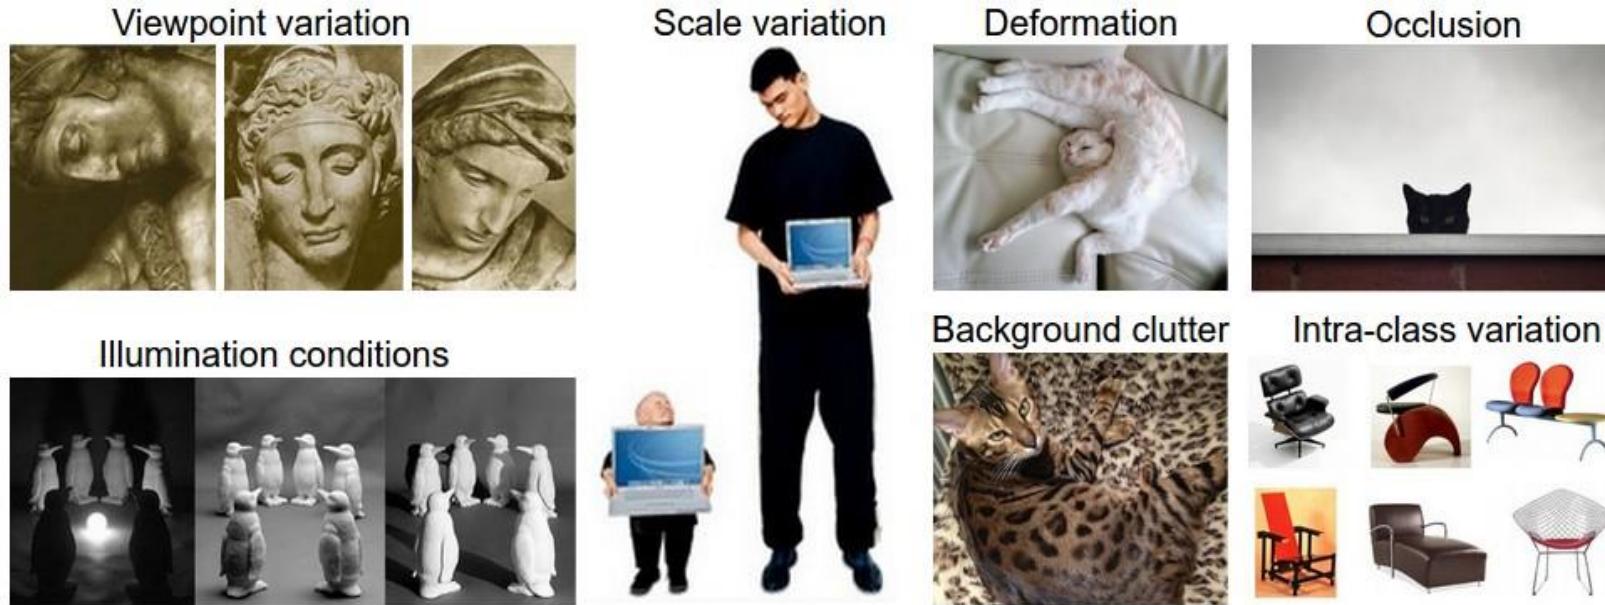
What computers see compared to what we see



401	402	403	404	405	406	407	408	409	410
411	412	413	414	415	416	417	418	419	420
421	422	423	424	425	426	427	428	429	430
431	432	433	434	435	436	437	438	439	440
441	442	443	444	445	446	447	448	449	450
451	452	453	454	455	456	457	458	459	460
461	462	463	464	465	466	467	468	469	470
471	472	473	474	475	476	477	478	479	480
481	482	483	484	485	486	487	488	489	490
491	492	493	494	495	496	497	498	499	500

# Classification Challenges

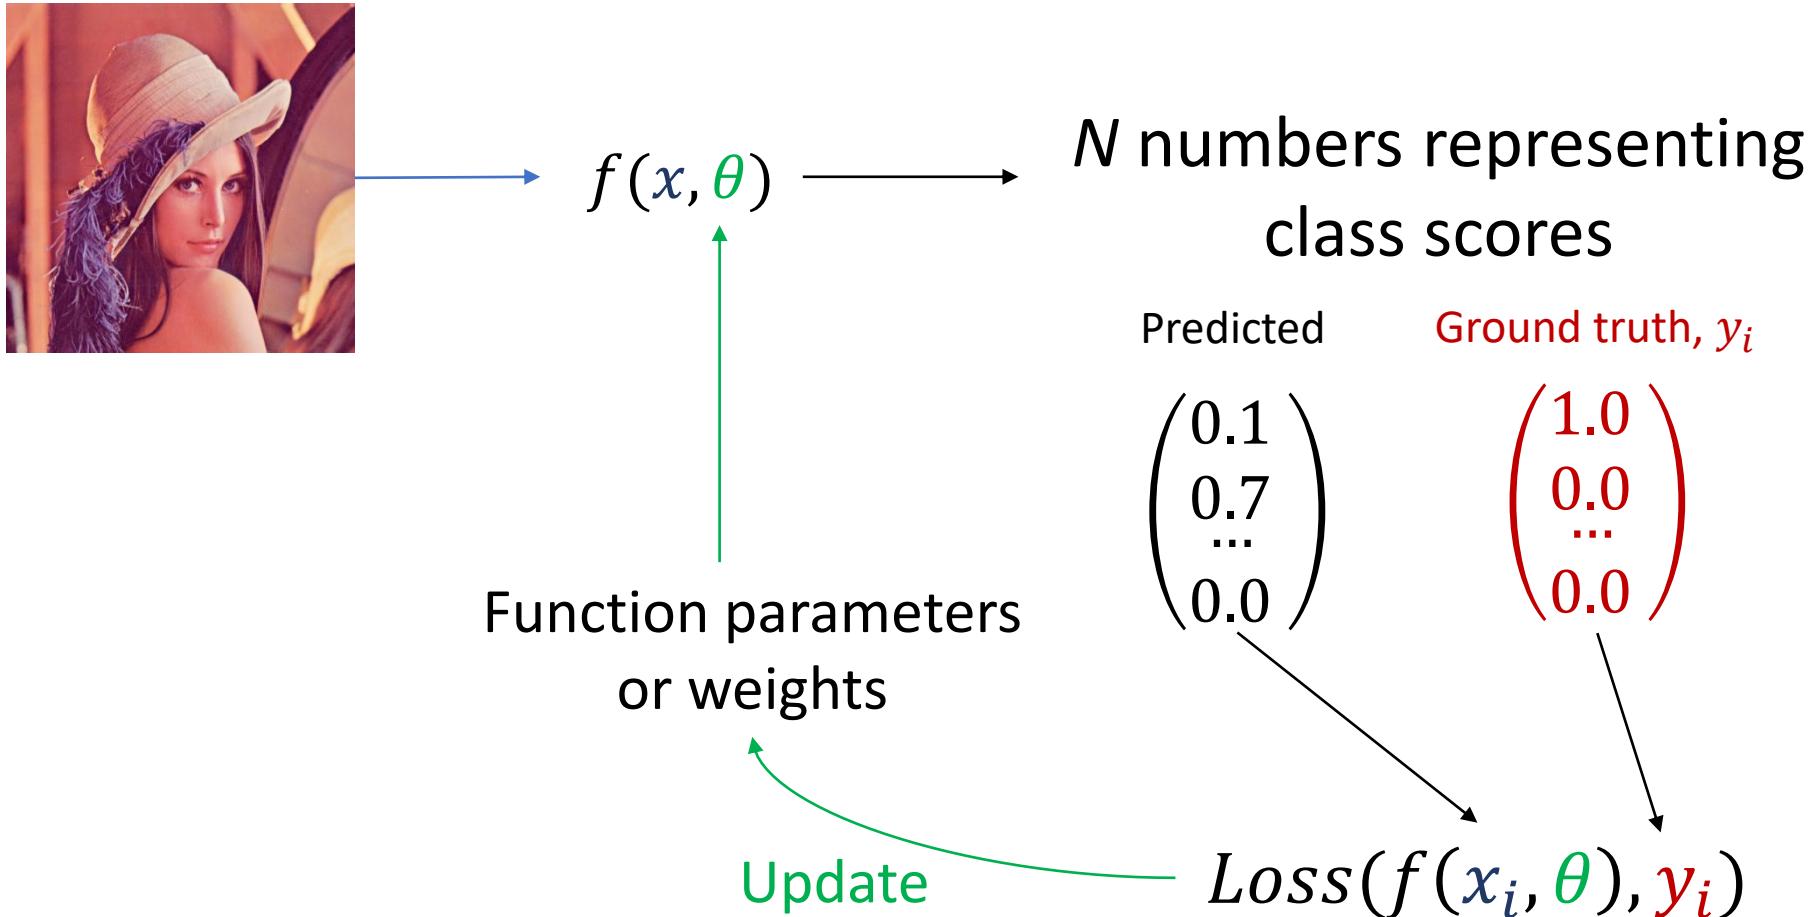
Directly specifying how a category looks like is impossible.



We need use a **Data Driven Approach**

# Supervised Learning

Find function  $f(x, \theta)$  that imitates a ground truth signal

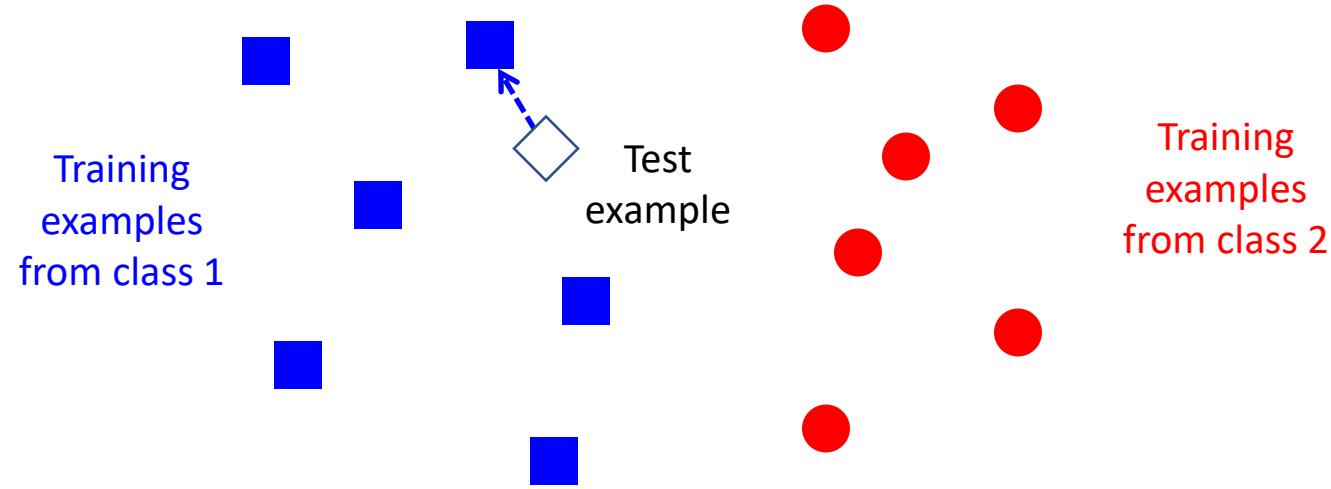


# Machine Learning Keywords

1. **Loss:** Quantify how good  $\theta$  are
2. **Optimization:** The process of finding  $\theta$  that minimize the loss
3. **Function:** Problem modelling → Deep networks are highly non-linear  $f(x, \theta)$

# Classifiers: K-Nearest neighbor

Features are represented in the descriptor space



$$f(\mathbf{x}, \boldsymbol{\theta}) = \text{label of the } K \text{ training examples nearest to } \mathbf{x}$$

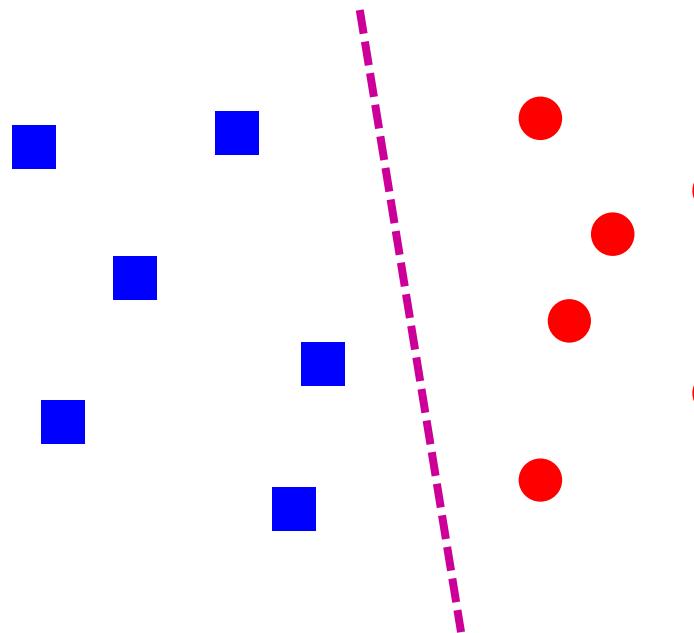
How fast is training? How fast is testing?

- $O(1), O(n)$

What is a good distance metric ? What K should be used? ☹

# Classifiers: Linear

Find a *linear function* to separate the classes:

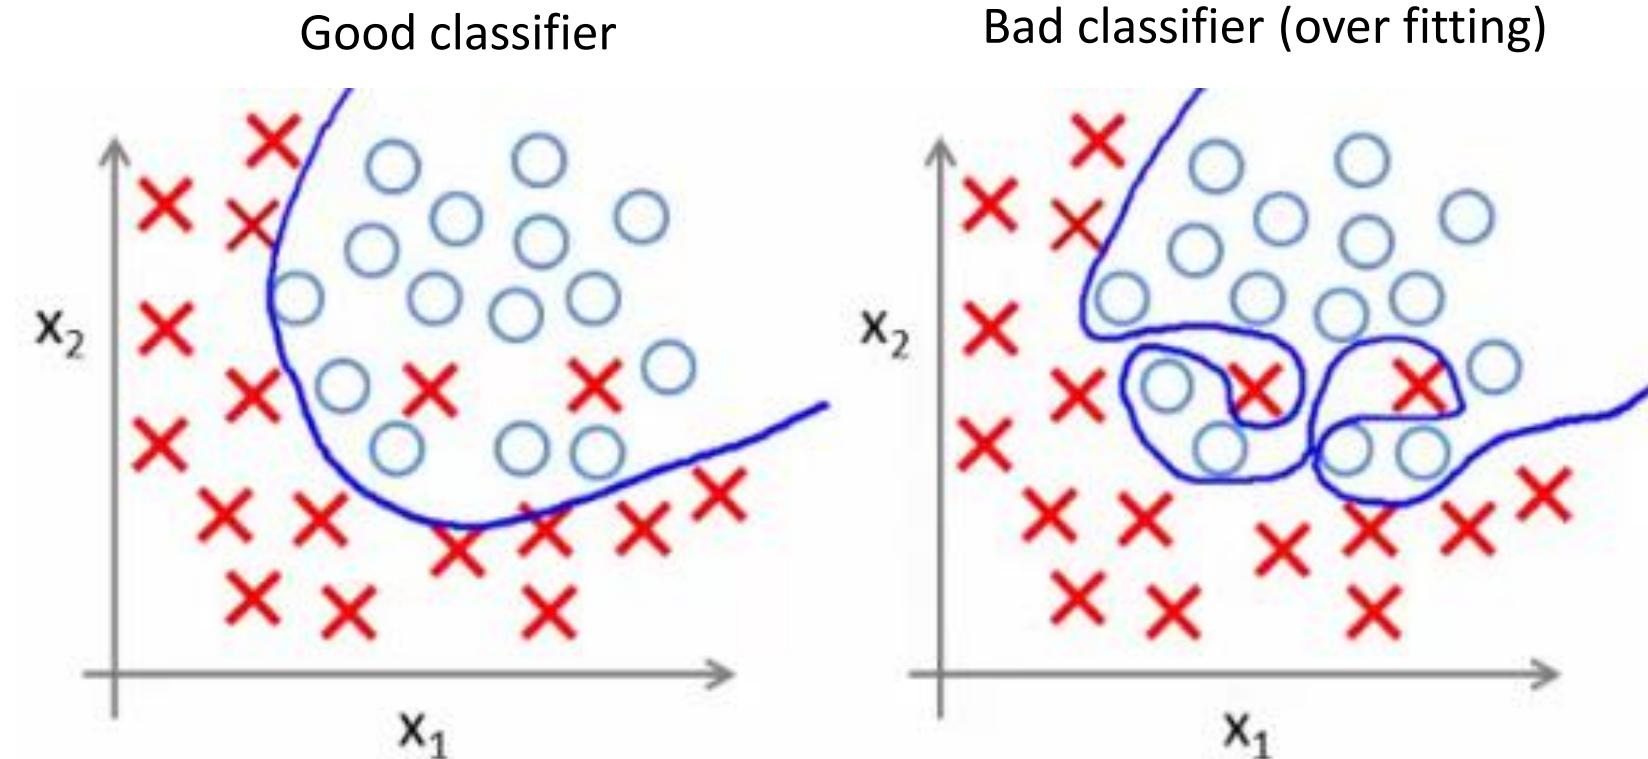


$$f(x, \theta) = \text{sgn}(\theta \cdot x + b)$$

What is  $\theta$ ? What is the dimensionality of images?

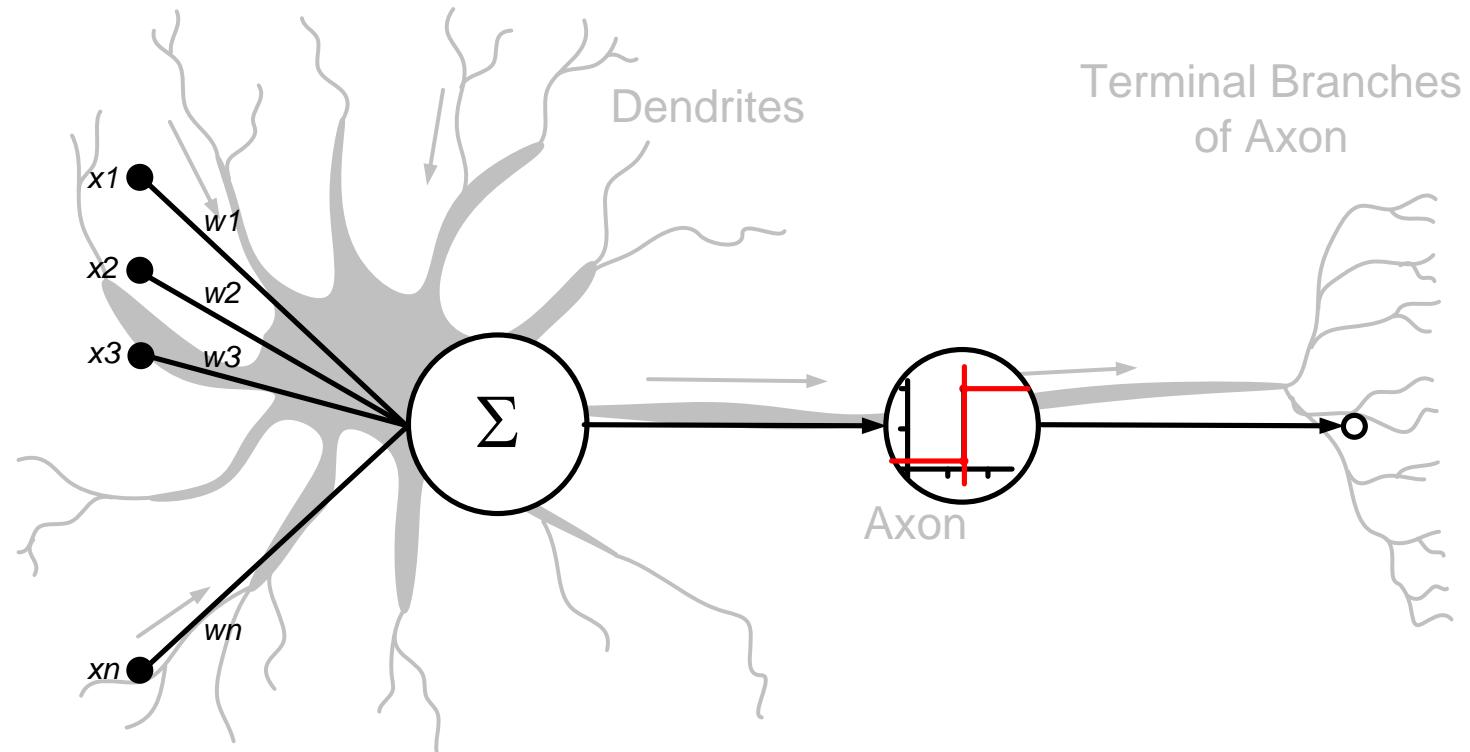
# Classifiers: non-linear

What is  $f(x, \theta)$  ?

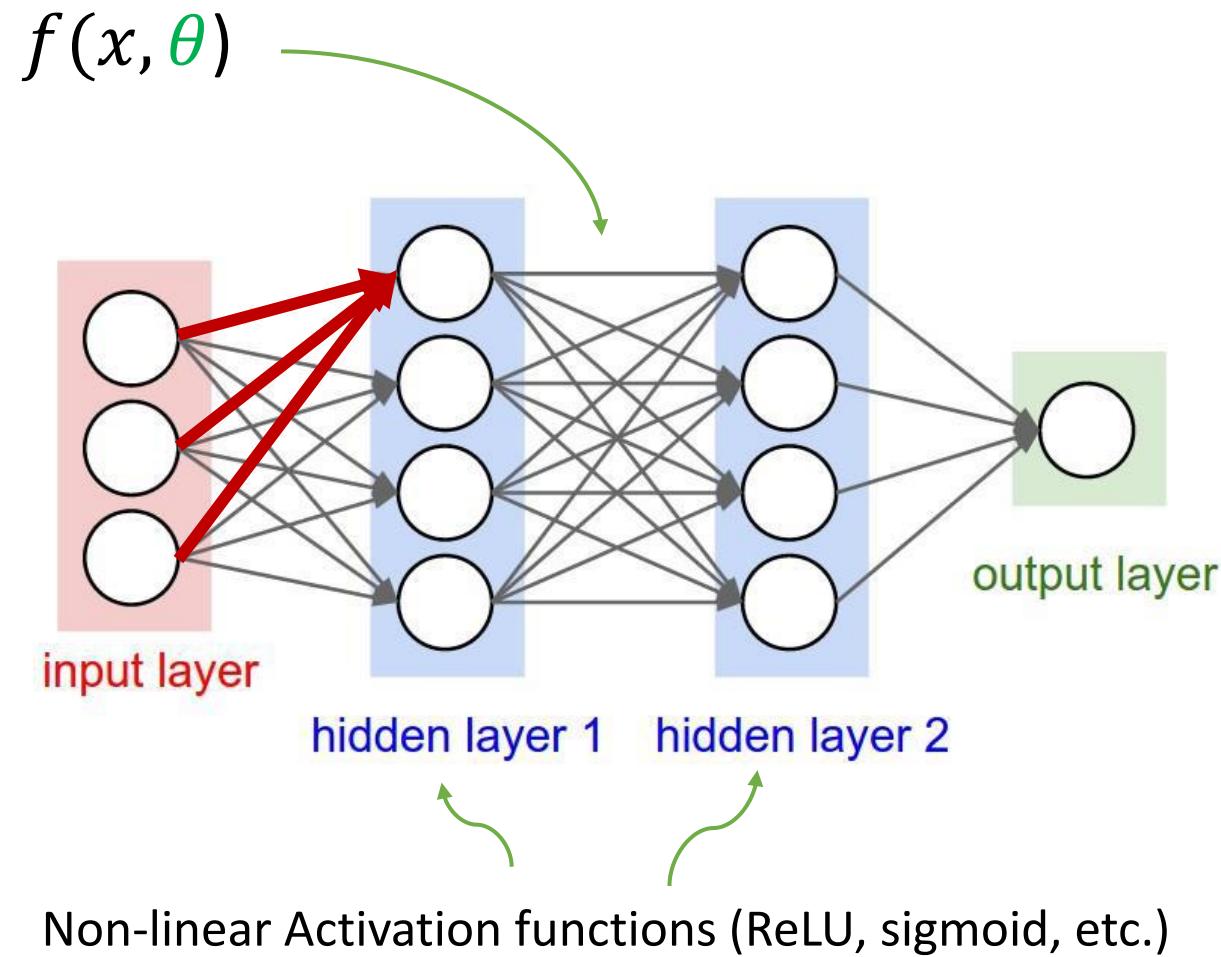


# Biological Inspiration

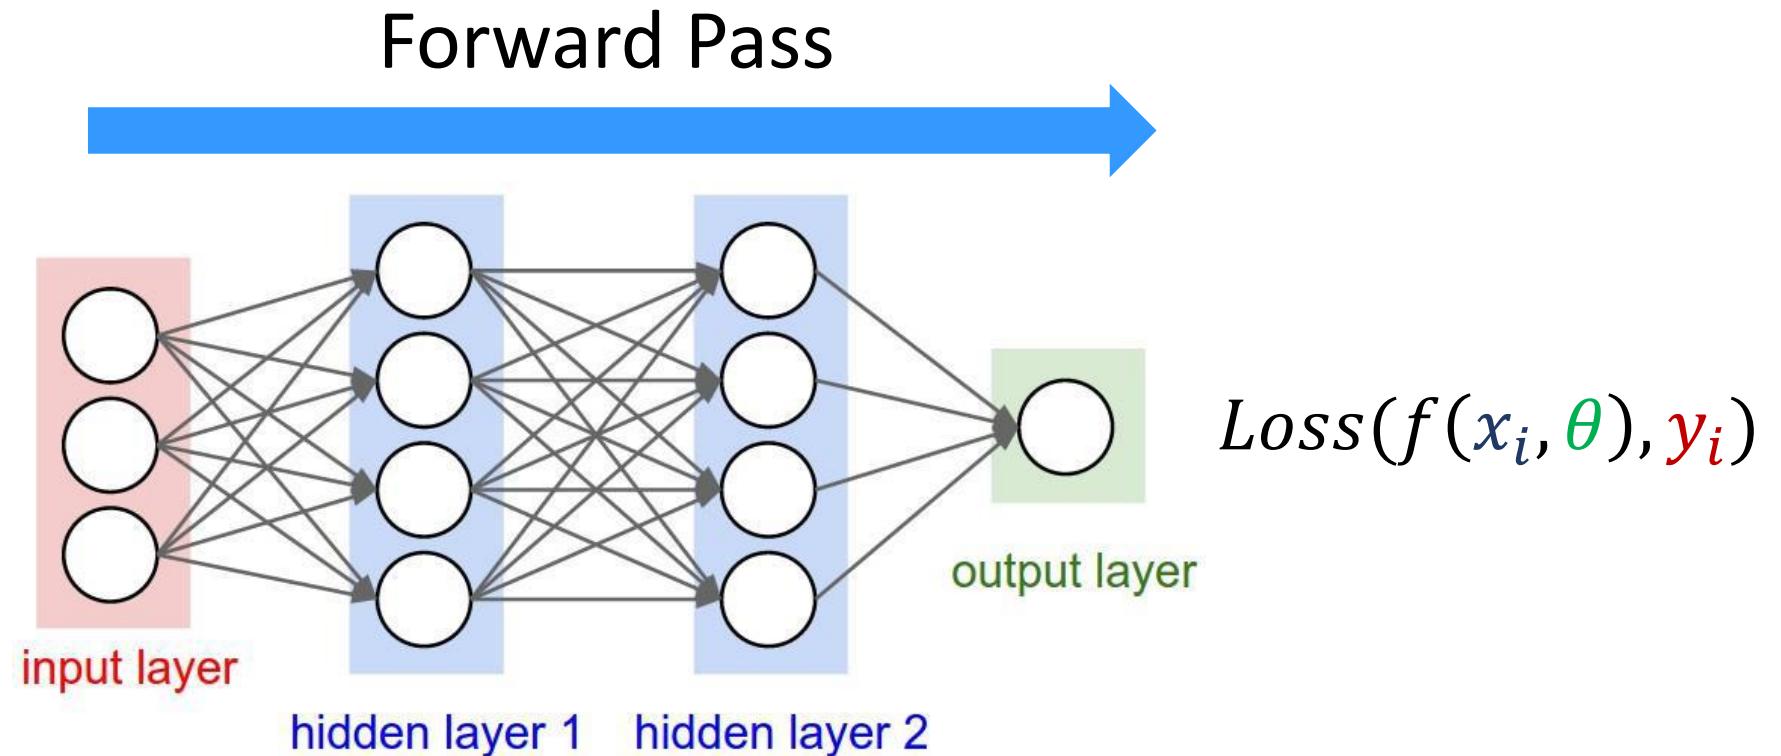
$f(x, \theta) = F(\theta x)$ ,  $F$  is a non-linear activation function (Step, ReLU, Sigmoid)



# Multi Layer Perceptron



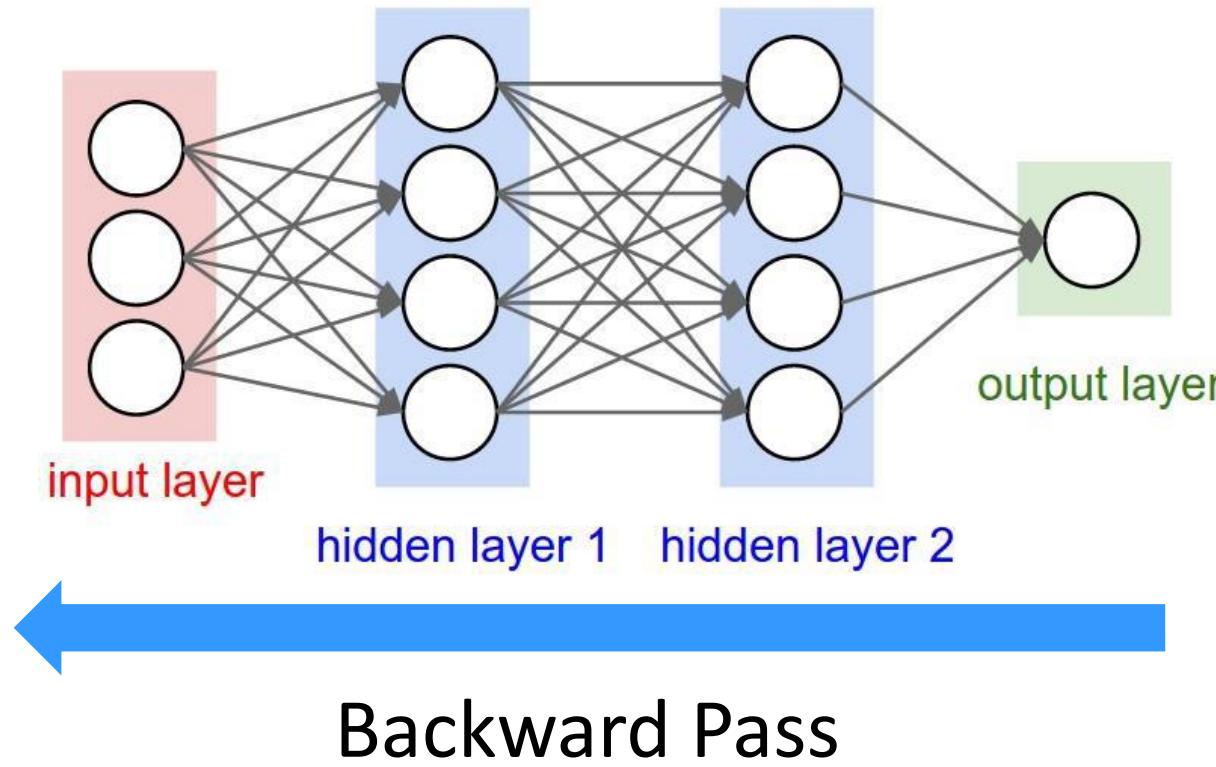
# Forward Propagation



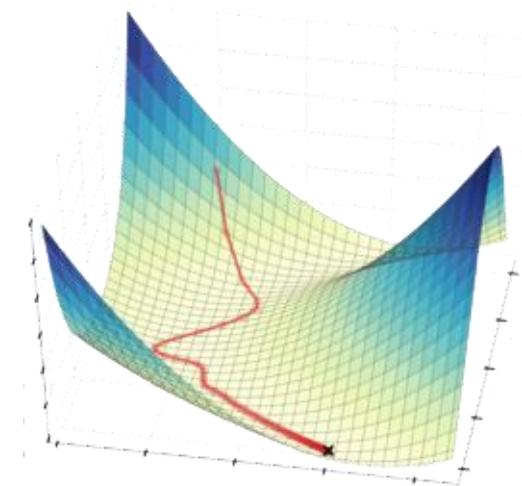
# Optimization: Back-propagation

Compute gradients with respect to all parameters and perform gradient descent

$$\theta_{new} = \theta_{old} - \mu \nabla_{\theta} Loss$$



$$Loss(f(x_i, \theta), y_i)$$

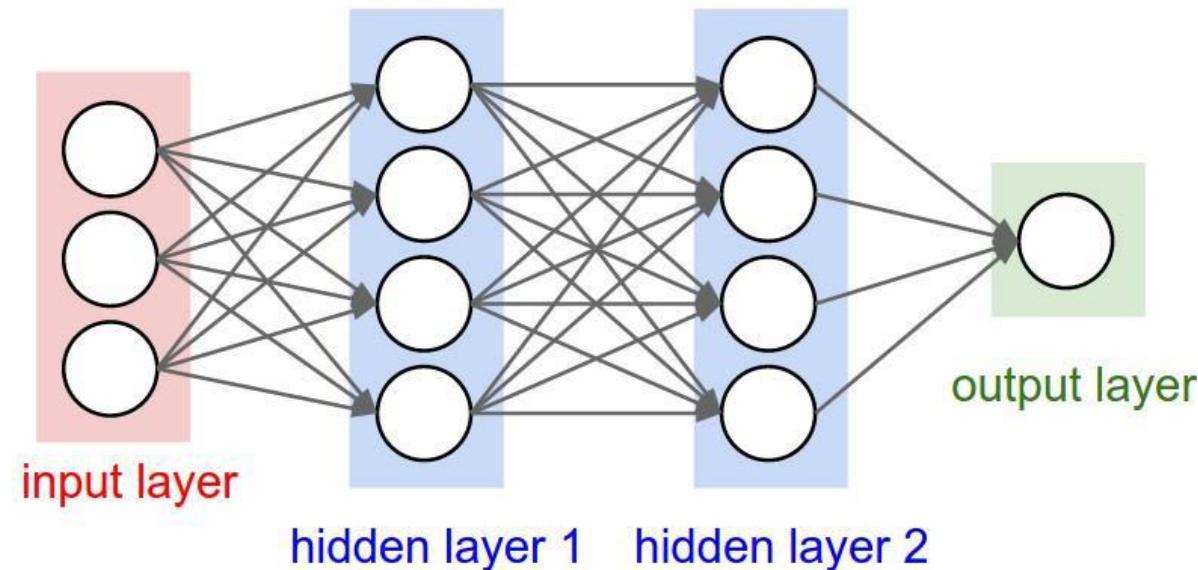


[1] Michael Nielsen, *Neural Networks and Deep Learning, Chapter 2* [PDF](#)

[2] Dreyfus, *Artificial Neural Networks, Back Propagation and the Kelley-Bryson Gradient Procedure*, Journal of Guidance, 1989. [PDF](#)

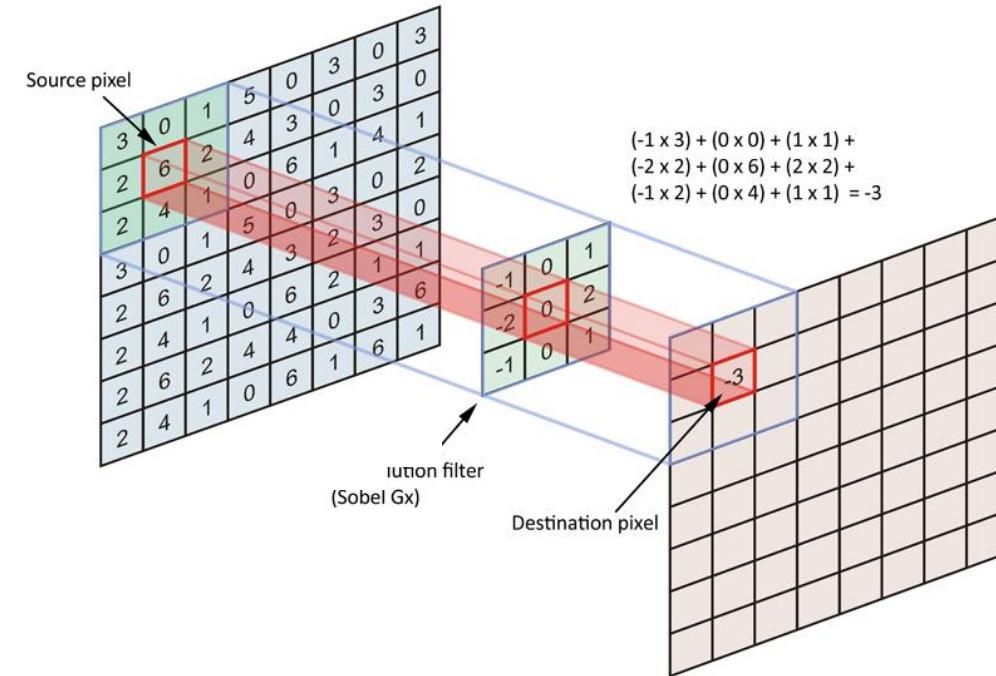
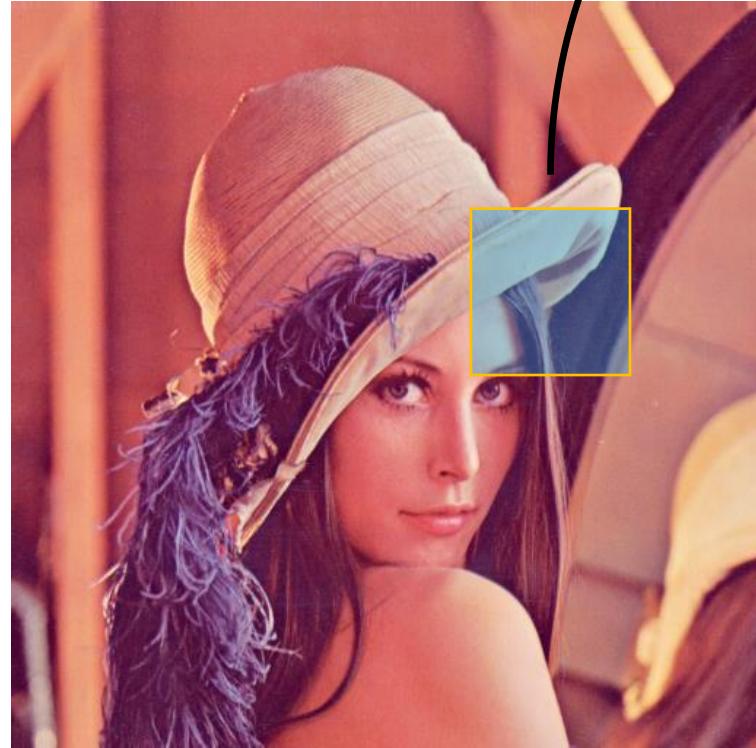
# Problems of fully connected network

Too many parameters → **possible overfitting.**

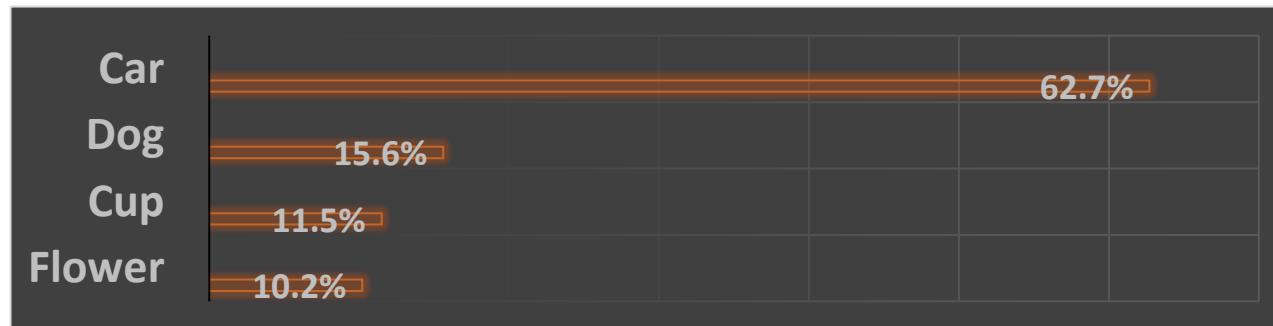
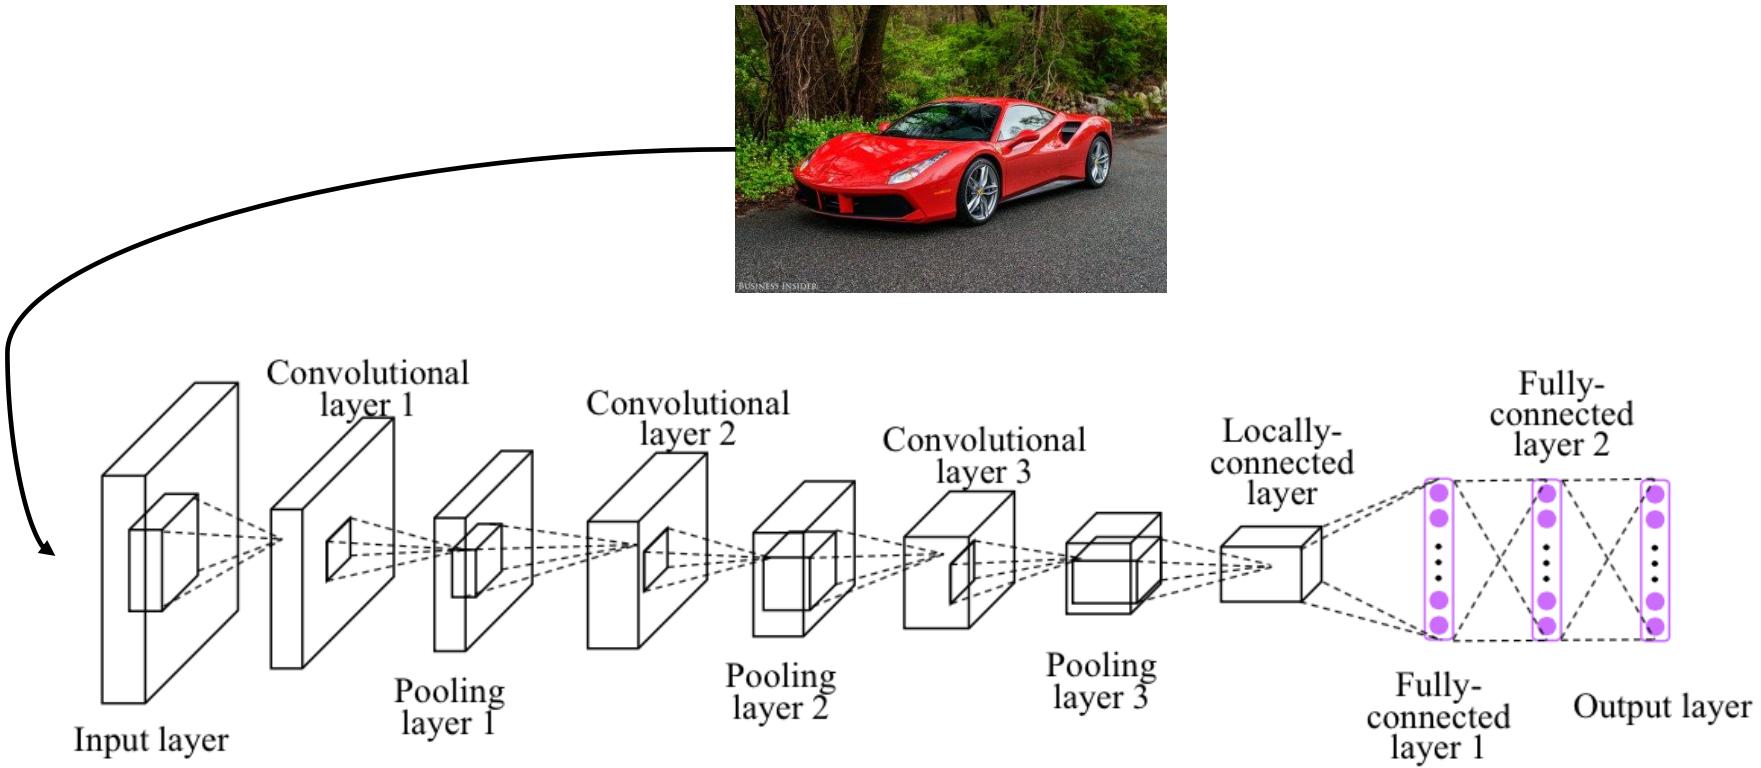


However, we are not using the fact that inputs are images!

# Convolutional Neural Networks

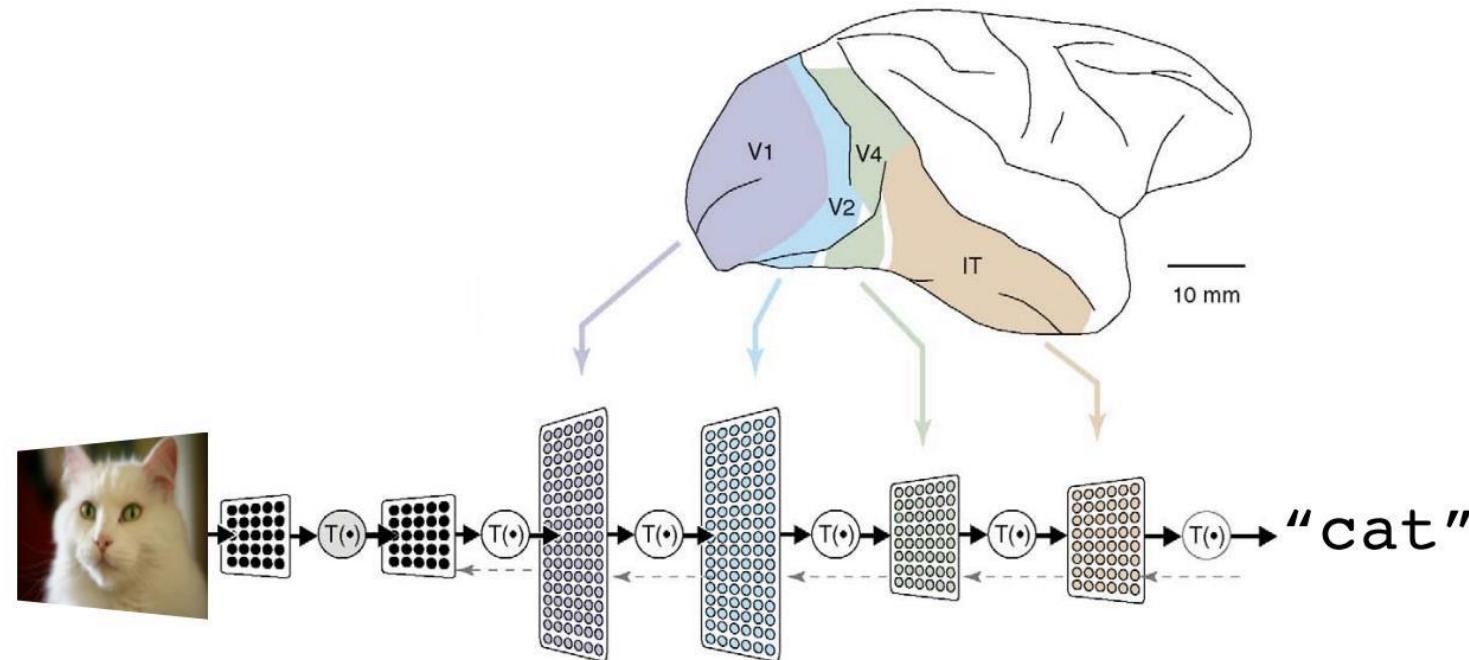


# Going Deep



# Why Deep?

1. Inspired by the **human visual system**
2. Learn **multiple layers** of transformations of input
3. Extract progressively more **sophisticated representations**

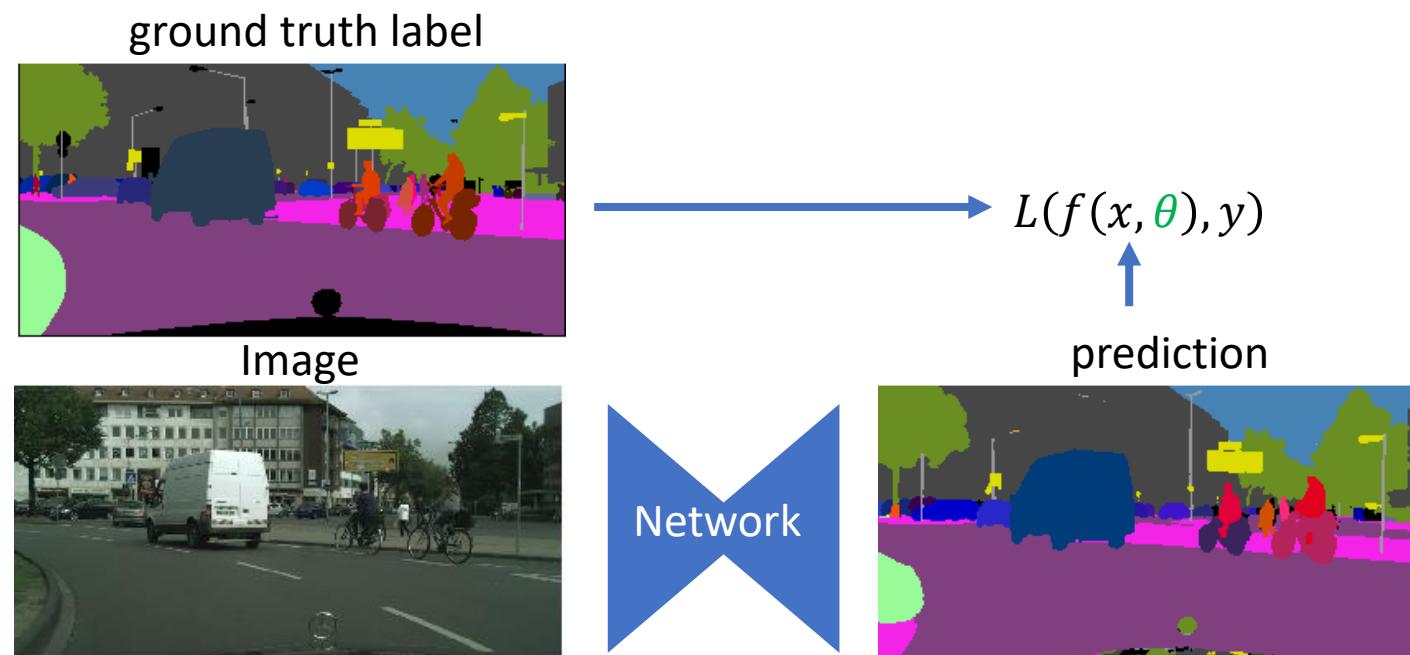


# Outline

- Introduction
- Supervised Learning
- Unsupervised Learning
- Applications to Computer Vision
- Conclusions
- Machine Learning for Drones

# Supervised Learning

- In supervised learning it is assumed that we have access to both input data or **images** and **ground truth labels**.
- Networks trained with supervision usually perform best
- However, data generation is hard, since it often must be **hand-labelled**



# Supervised Learning

- Image Segmentation



Long, Shelhamer, *Fully Convolutional Networks for Semantic Segmentation*,  
Conference of Computer Vision and Pattern Recognition (CVPR), 2015. [PDF](#)

# Supervised Learning

- Image Captioning



"little girl is eating piece of cake."



"baseball player is throwing ball in game."



"woman is holding bunch of bananas."



"black cat is sitting on top of suitcase."



"a young boy is holding a baseball bat."



"a cat is sitting on a couch with a remote control."



"a woman holding a teddy bear in front of a mirror."



"a horse is standing in the middle of a road."

# Outline

- Introduction
- Supervised Learning
- Unsupervised Learning
- Applications to Computer Vision
- Conclusions
- Machine Learning for Drones

# Unsupervised Learning

- In unsupervised learning we only have access to input data or **images**.
- Usually, these methods are more popular because they can use much larger datasets that do not need to be manually labelled.



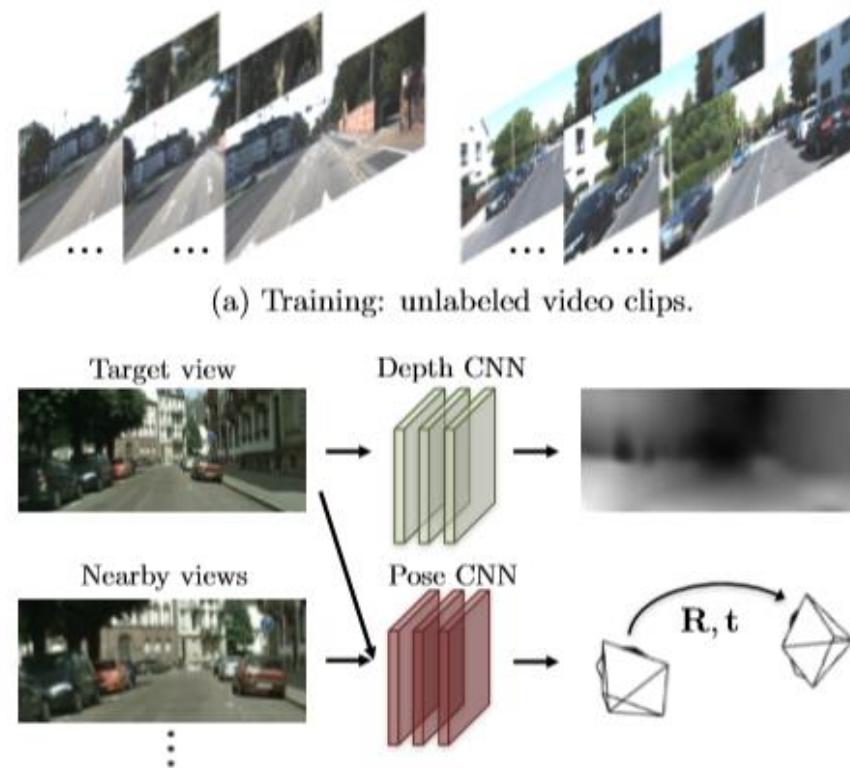
# Unsupervised Learning

- Monocular Depth Estimation



# Unsupervised Learning

- Structure from Motion



# Unsupervised vs. Supervised learning

	<b>Supervised</b>	<b>Unsupervised</b>
Performance	Usually better for the same dataset size.	Usually worse, but can outperform supervised methods due to larger data availability.
Data availability	Low, due to manual labelling.	High, no labelling required.
Training	Simple, ground truth gives a strong supervision signal.	Sometimes difficult, loss functions have to be engineered to get good results.
Generalizability	Good, although sometimes the network learns to blindly copy the labels provided, leading to poor generalizability.	Better, since unsupervised losses often encode the task in a more fundamental way.

# Outline

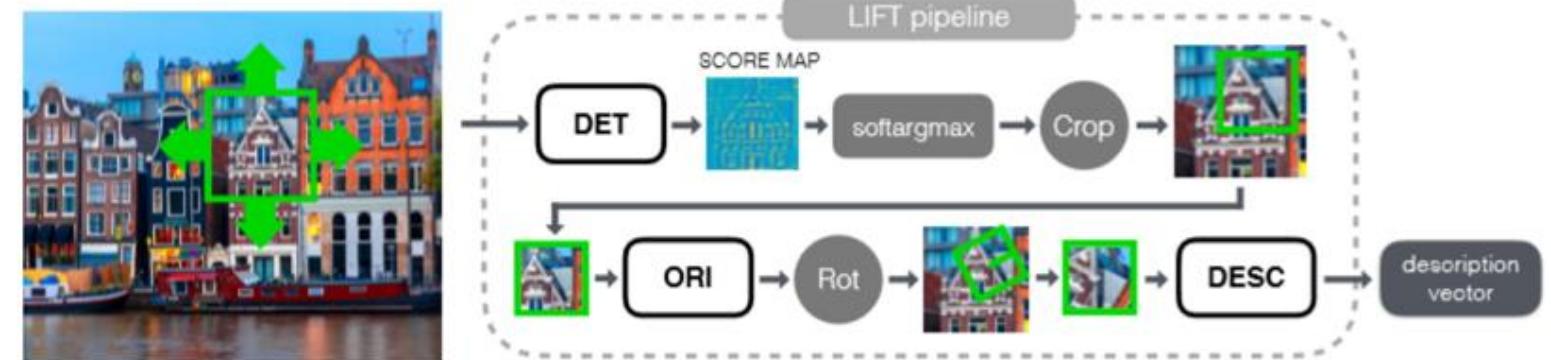
- Introduction
- Supervised Learning
- Unsupervised Learning
- Applications to Computer Vision
- Conclusions
- Machine Learning for Drones

# Keypoint Detection and Description

- Deep Descriptors: Learned Invariant Feature Transform (LIFT)

LIFT consists of 3 neural networks:

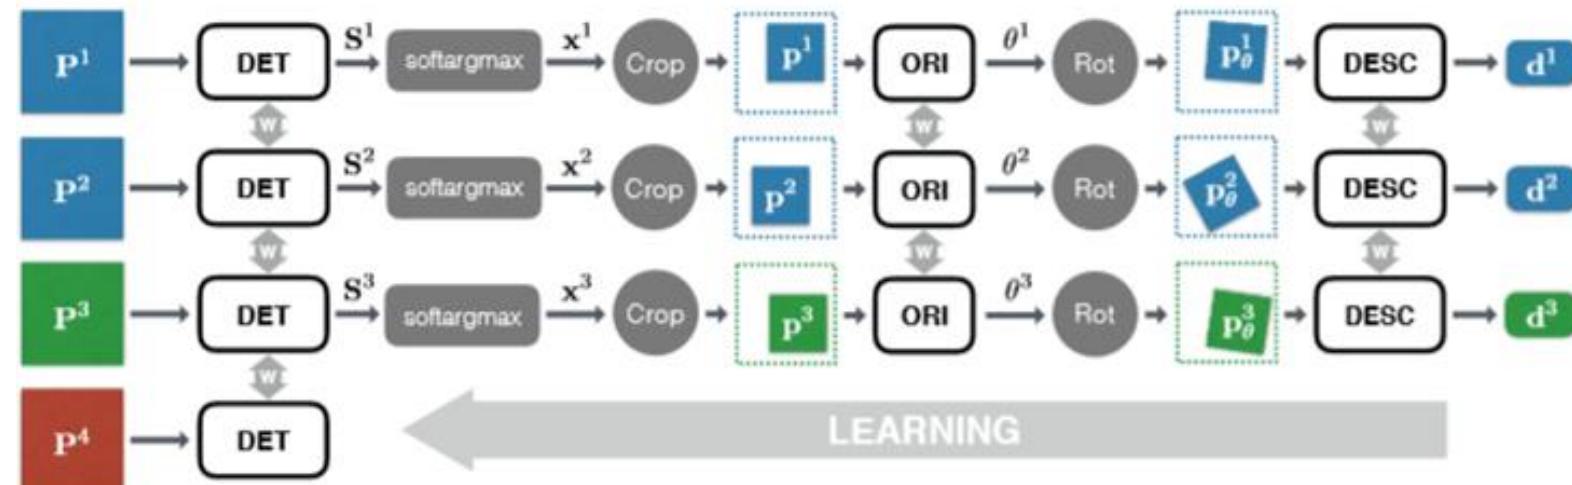
- A keypoint detector
- An orientation detector
- A descriptor generator



# LIFT Loss

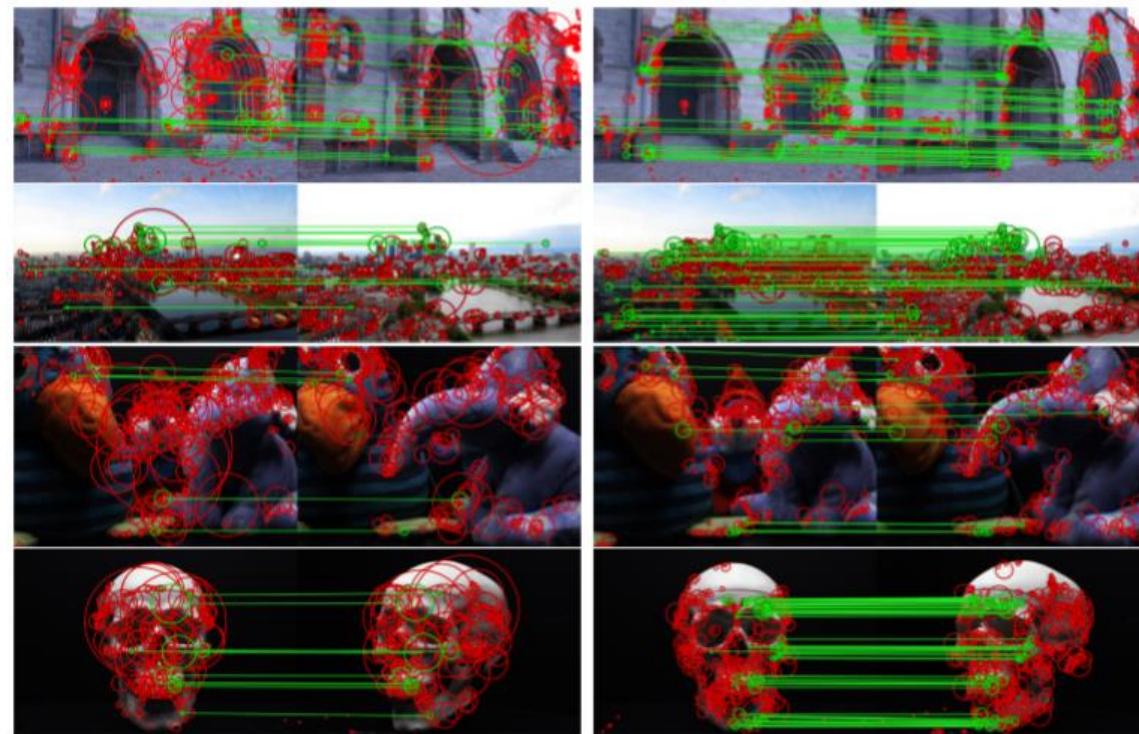
The LIFT loss has 3 components:

- Descriptors of **matching patches**,  $d^1, d^2$ , should be **close**
- Descriptors of **non-matching patches**,  $d^1, d^2$ , should be **far apart**
- Keypoints should not be in homogeneous regions: P4 should not be detected as a keypoint.



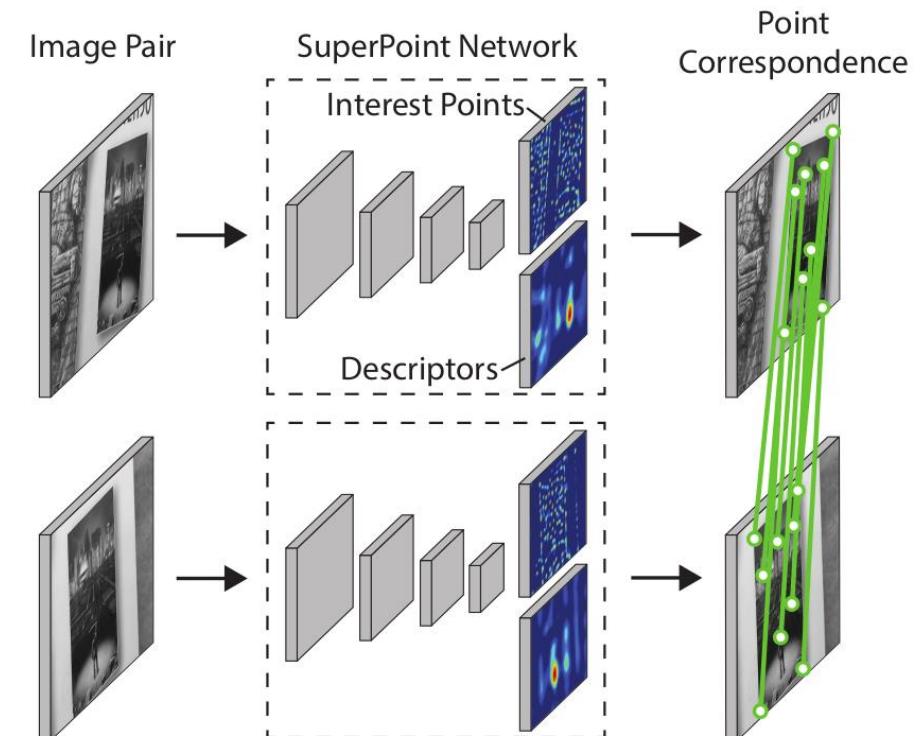
# LIFT - Results

- Works better than SIFT! (well, in some datasets)



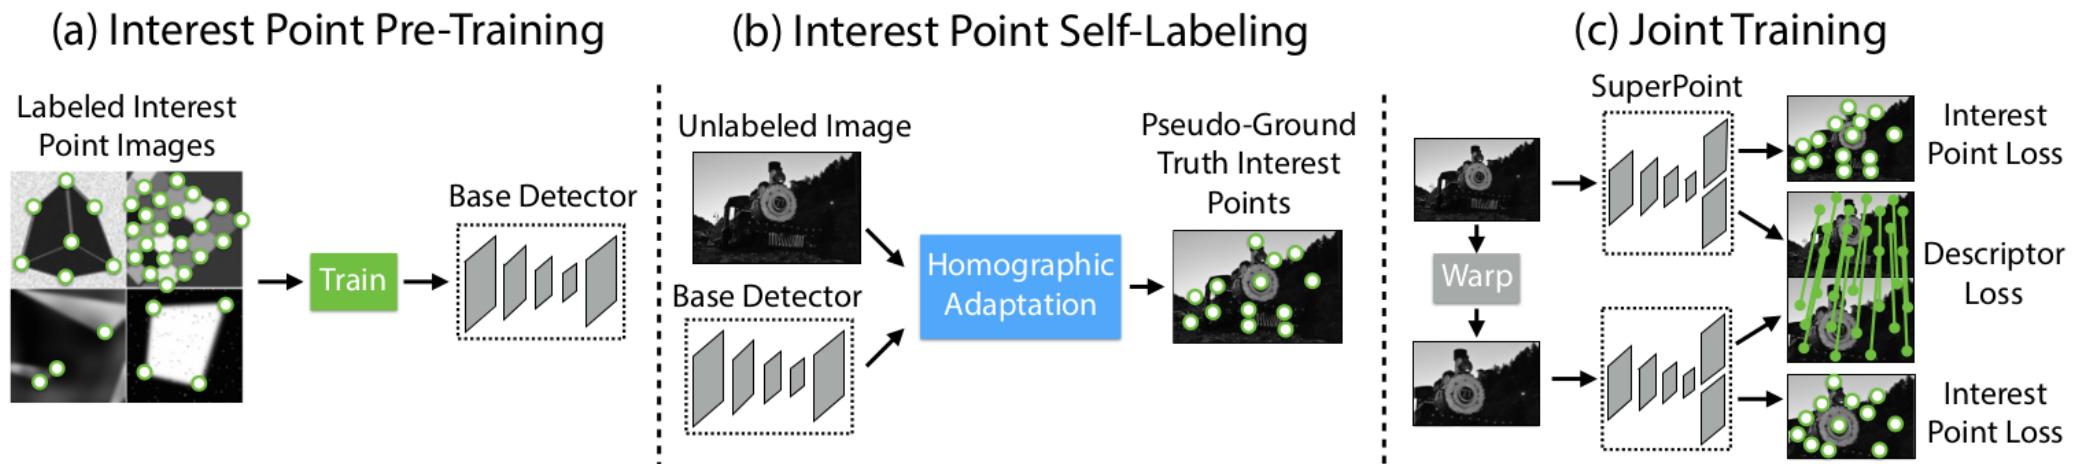
# Keypoint Detection and Description

- SuperPoint: Self-Supervised Interest Point Detection and Description
  - SIFT and friends are complicated heuristic algorithms
  - Still, SIFT is a hard to beat baseline for new methods
  - Can we do better with a data driven approach?
  - SuperPoint uses a CNN to **predict keypoints and descriptors simultaneously.**
  - Detector less accurate than SIFT, but descriptor shown to outperform SIFT in some scenarios.

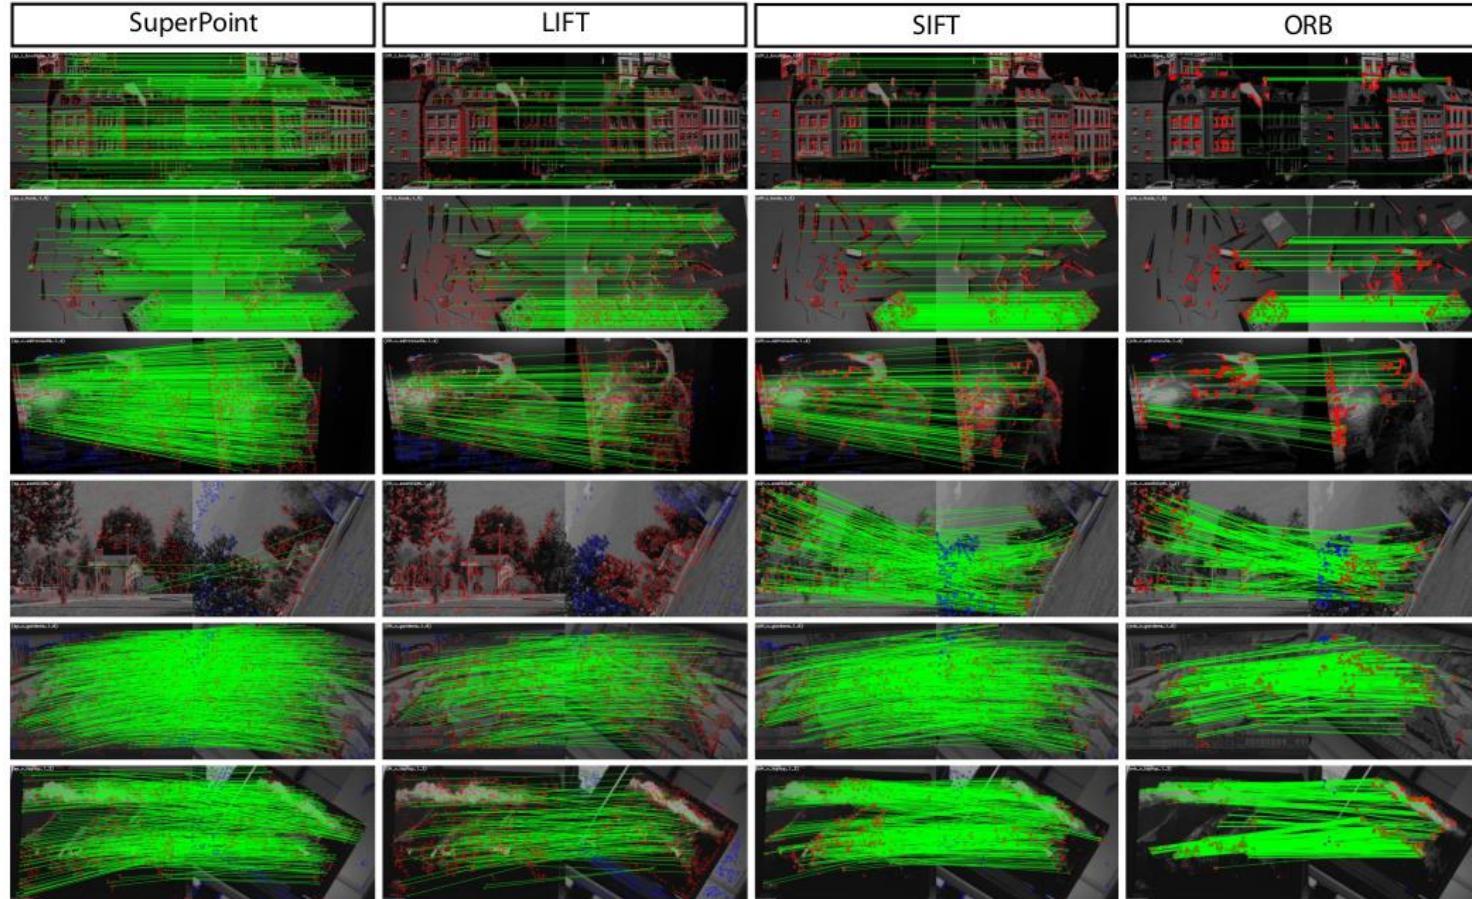


# SuperPoint - Training

- Training Steps:
  - Training on synthetic dataset to bootstrap the detector
  - Use this detector on real images and generate groundtruth correspondences by sampling perspective transformations
  - Jointly training both detector and descriptor



# SuperPoint - Results



# SuperPoint - Results

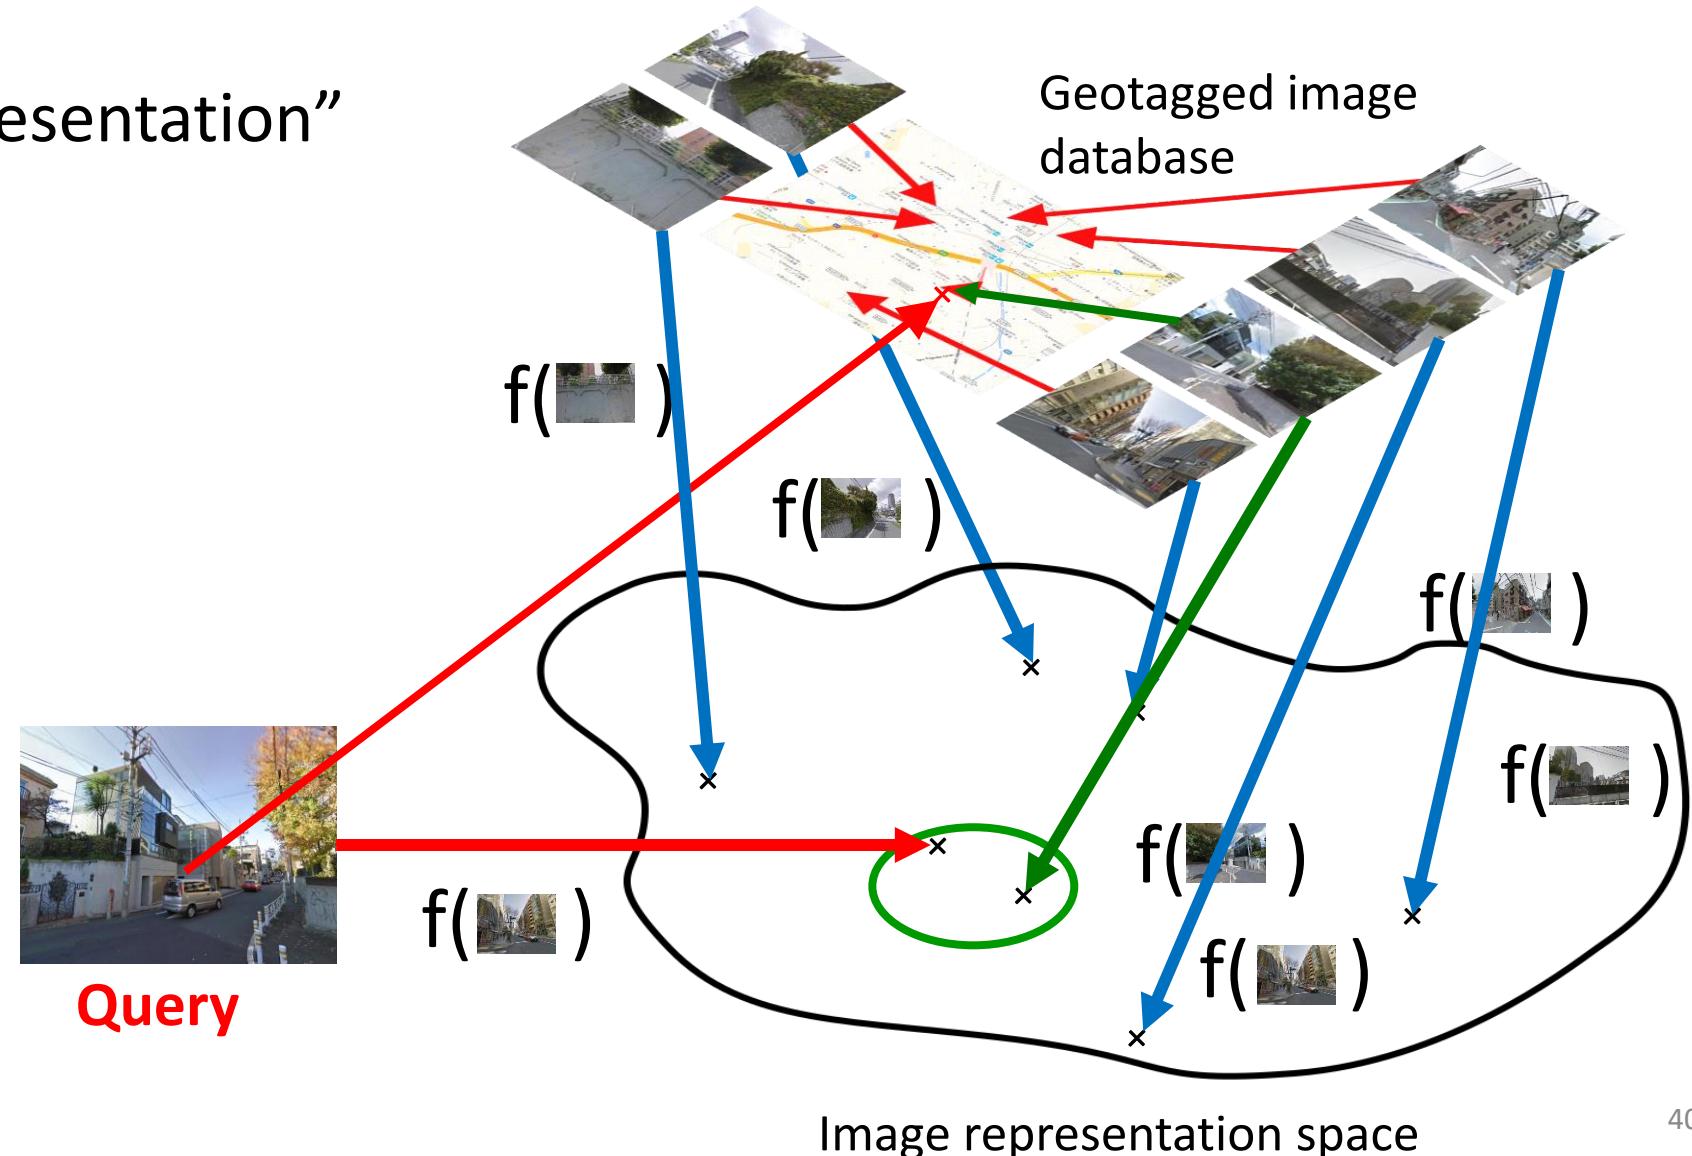
	Homography Estimation			Detector Metrics		Descriptor Metrics	
	$\epsilon = 1$	$\epsilon = 3$	$\epsilon = 5$	Rep.	MLE	NN mAP	M. Score
<i>SuperPoint</i>	.310	<b>.684</b>	<b>.829</b>	.581	1.158	<b>.821</b>	<b>.470</b>
<i>LIFT</i>	.284	.598	.717	.449	1.102	.664	.315
<i>SIFT</i>	<b>.424</b>	.676	.759	.495	<b>0.833</b>	.694	.313
<i>ORB</i>	.150	.395	.538	<b>.641</b>	1.157	.735	.266

## Legend

- **$\epsilon$** : Pixel distance to count keypoint as correct. SIFT is very accurate due to sub-pixel refinement.
- **MLE** (Mean Localization Error): Lower is better. Metric for determining detector accuracy
- **NN mAP** (nearest-neighbor mean average precision): Measures discriminativeness of descriptors. Higher is better.
- **M. Score** (Matching Score): Evaluates detector and descriptor jointly on groundtruth correspondences.

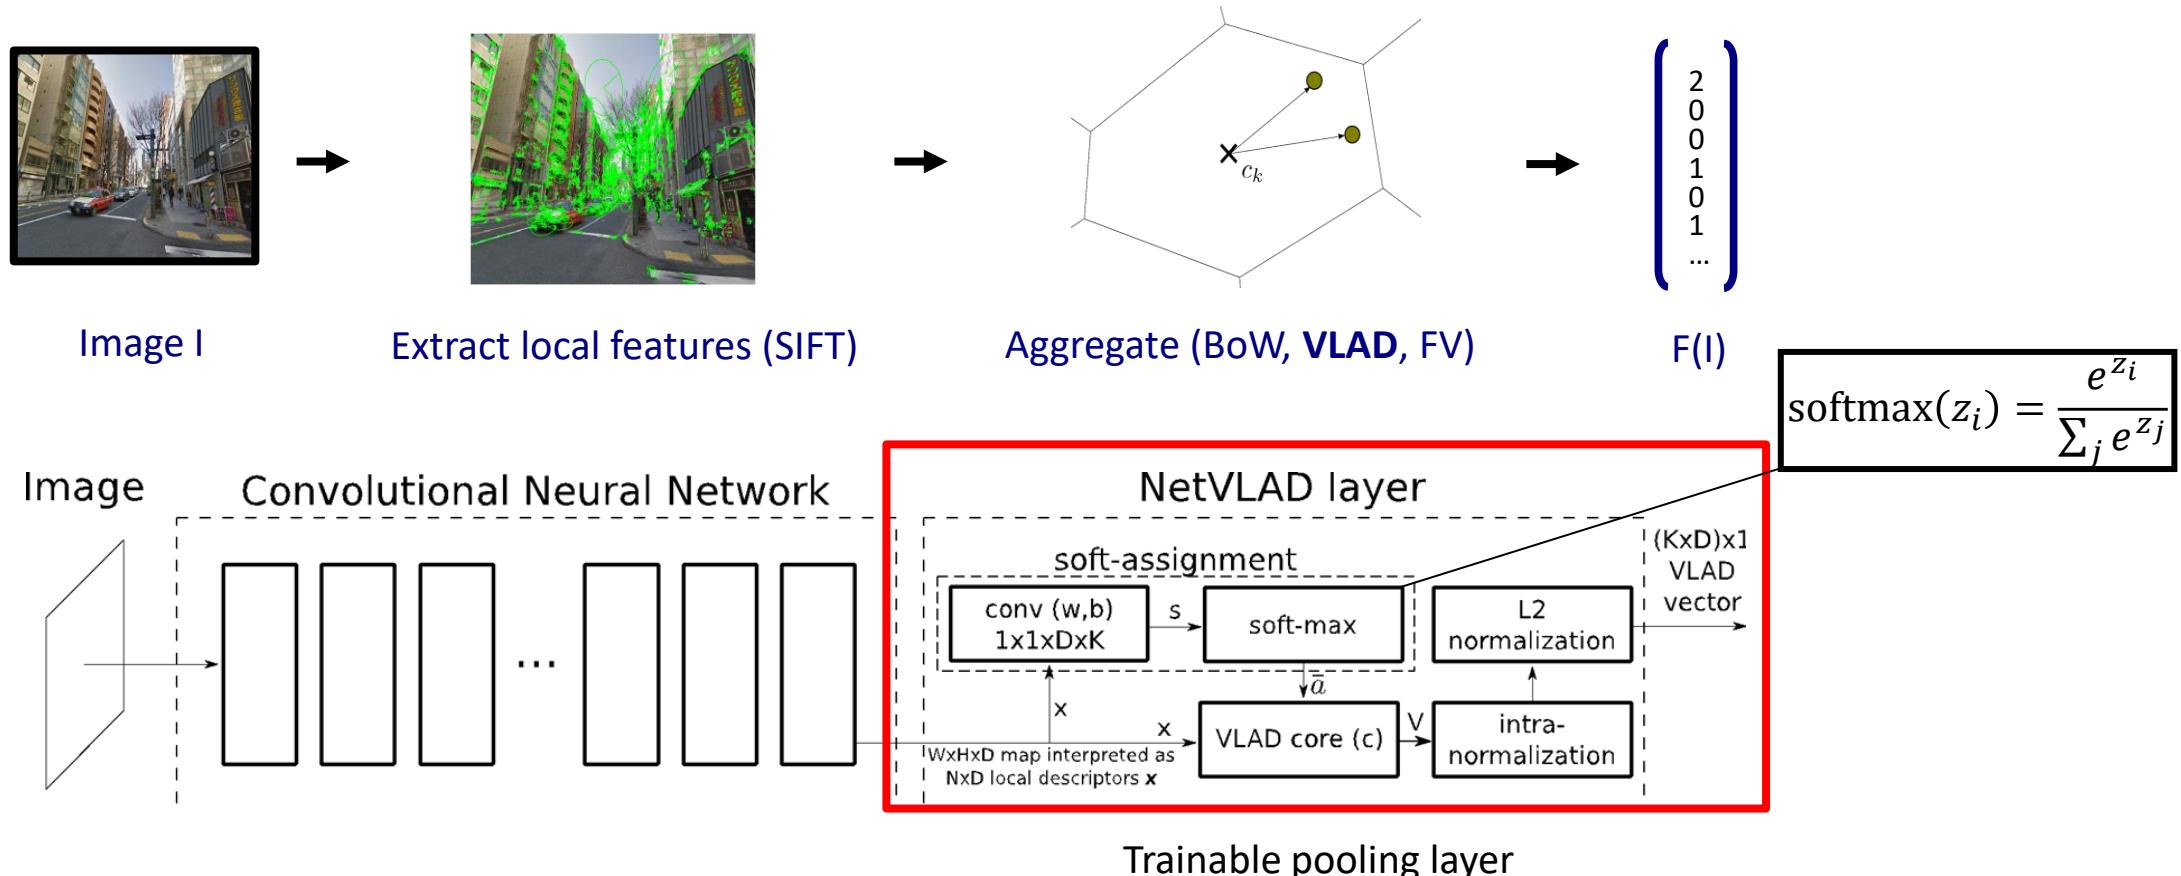
# Place Recognition

- Design an “image representation” extractor  $f(I, \theta)$



# NetVLAD

- Mimic the classical pipeline with deep learning



# NetVLAD - Loss

- Triplet loss formulation

$$D_p = ||F_\theta(\text{[img]}) - F_\theta(\text{[img]})||^2 \rightarrow \text{Matching samples}$$

$$D_n = ||F_\theta(\text{[img]}) - F_\theta(\text{[img]})||^2 \rightarrow \text{Non matching samples}$$

$$L_\theta = \sum_{samples} \max(D_{p(\theta)} + m - D_{n(\theta)}, 0)$$

margin  
↑

Disclaimer: The actual NetVlad loss is a slightly more complicated version of the one above

# NetVLAD - Results

- Code, dataset and trained network online: give it a try [here!](#)

Query



Top result

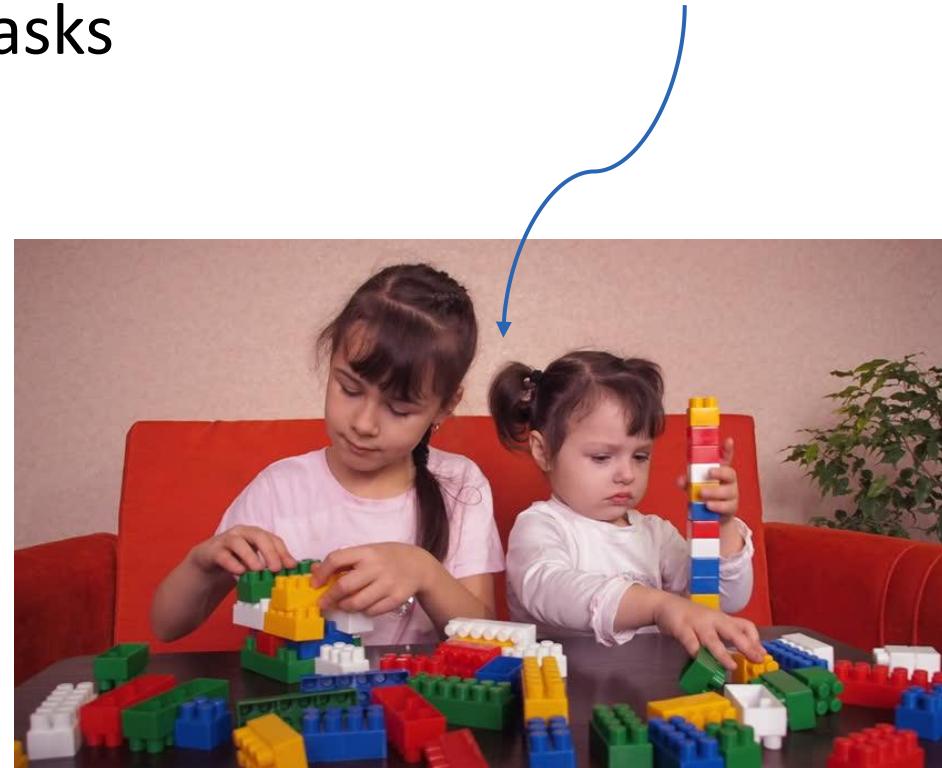
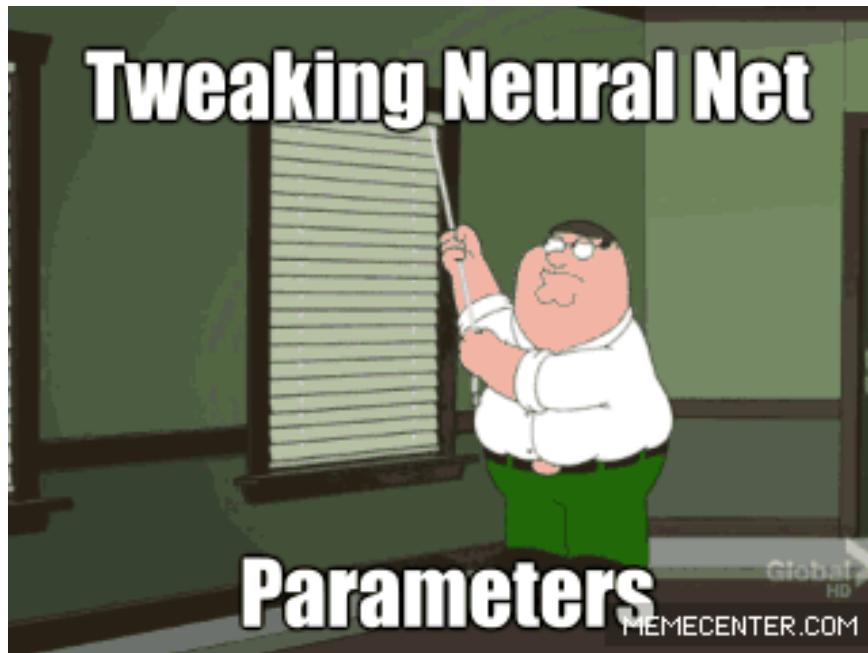


Green: Correct   Red: Incorrect

# Deep Learning Limitations

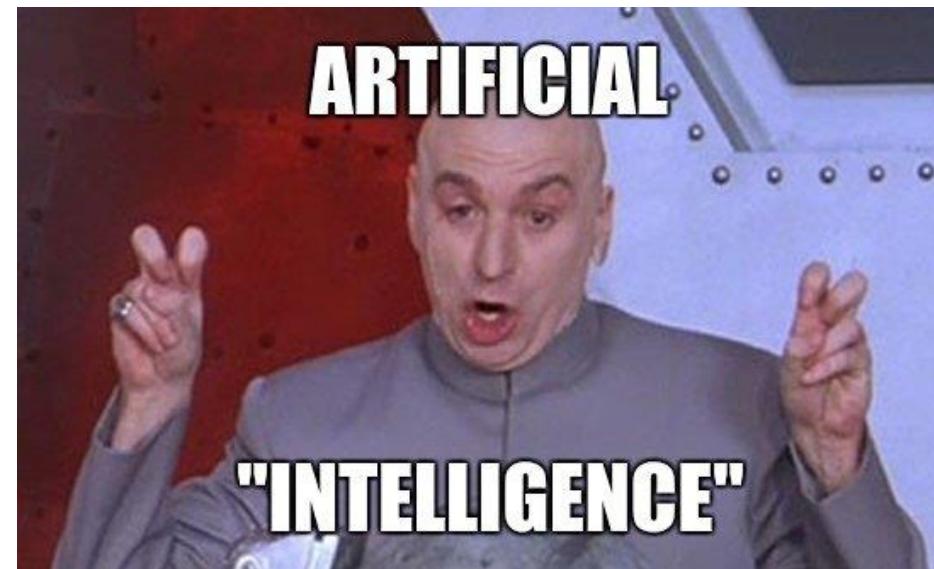
- Requires a **lot of data** to learn
- **Difficult debugging** and finetuning
- **Poor generalization** across similar tasks

**Neural Networks  
Practitioners**



# Things to Remember

- Deep Learning is able to **extract meaningful patterns** from data.
- It can be applied to a **wide range of tasks**.
- **Artificial Intelligence ≠ Deep Learning**



# Come over for projects in DL!

- Visit our webpage for projects! [http://rpg.ifi.uzh.ch/student\\_projects.php](http://rpg.ifi.uzh.ch/student_projects.php)

## Asynchronous Processing for Event-based Deep Learning - Available



**Description:** Event cameras such as the Dynamic Vision Sensor (DVS) are recent sensors with large potential for high-speed and high dynamic range robotic applications. Since their output is sparse traditional algorithms, which are designed for dense inputs such as frames, are not well suited. The goal of this project is explore ways to adapt existing deep learning algorithms to handle sparse asynchronous data from events. Applicants should have experience in C++ and python deep learning frameworks (tensorflow or pytorch), and have a strong background in computer vision.

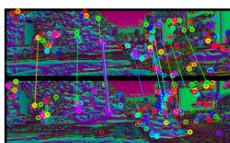
**Goal:** The goal of this project is explore ways to adapt existing deep learning algorithms to handle sparse asynchronous data from events.

**Contact Details:** Daniel Gehrig (dgehrig at ifi.uzh.ch)

**Thesis Type:** Semester Project / Master Thesis

[See project on SIROP](#)

## Data-driven Keypoint Extractor for Event Data - Available



**Description:** Neuromorphic cameras exhibit several amazing properties such as robustness to HDR scenes, high-temporal resolution, and low power consumption. Thanks to these characteristics, event cameras are applied for camera pose estimation for fast motions in challenging scenes. A common technique for camera pose estimation is the extraction and tracking of keypoints on the camera plane. In the case of event cameras, most existing keypoint extraction methods are handcrafted manually. As a new promising direction, this project tackles the keypoint extraction in a data-driven fashion based on recent advances in frame-based keypoint extractors.

**Goal:** The project aims to develop a data-driven keypoint extractor, which computes interest points in event data. Based on the current advances of learned keypoint extractors for traditional frames, the approach will leverage neural network architectures to extract and describe keypoints in an event stream. The student should have prior programming experience in a deep learning framework and completed at least one course in computer vision.

**Contact Details:** Contact Details: Nico Messikommer [nmessi (at) ifi (dot) uzh (dot) ch], Mathias Gehrig [mgehrig (at) ifi (dot) uzh (dot) ch]

**Thesis Type:** Semester Project / Master Thesis

[See project on SIROP](#)

## Designing a New Event Camera with Events and Images - Available



**Description:** Event cameras such as the Dynamic Vision Sensor (DVS) are recent sensors with a lot of potential for high-speed and high dynamic range robotic applications. They have been successfully applied in many applications, such as high speed video and high speed visual odometry. Due to their high speed and

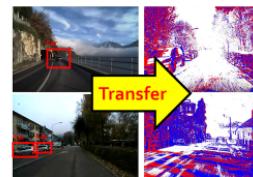
**Goal:** The goal of this project is to design a new event camera that combines events and standard images.

**Contact Details:** Daniel Gehrig (dgehrig (at) ifi (dot) uzh (dot) ch), Mathias Gehrig (mgehrig (at) ifi (dot) uzh (dot) ch)

**Thesis Type:** Semester Project / Master Thesis

[See project on SIROP](#)

## Domain Transfer between Events and Frames - Available



**Description:** During the last years, a vast collection of frame-based datasets was collected for countless tasks. In comparison, event-based datasets represent only a tiny fraction of the available datasets. Thus, it is highly promising to use labelled frame datasets to train event-based networks as current data-driven approaches heavily rely on labelled data.

**Goal:** In this project, the student extends current advances from the UDA literature for traditional frames to event data in order to transfer multiple tasks from frames to events. The approach should be validated on several tasks (segmentation, object detection, etc.) in challenging environments (night, high-dynamic scenes) to highlight the benefits of event cameras. As several deep learning methods are used as tools for the task transfer, a strong background in deep learning is required. If you are interested, we are happy to provide more details.

**Contact Details:** Nico Messikommer [nmessi (at) ifi (dot) uzh (dot) ch], Daniel Gehrig [dgehrig (at) ifi (dot) uzh (dot) ch]

**Thesis Type:** Semester Project / Master Thesis

[See project on SIROP](#)

# Additional Readings

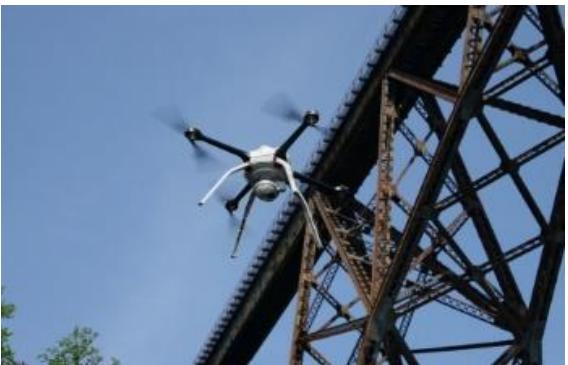
- Nielsen, *Neural Networks and Deep Learning*, 2018. [PDF](#)
- Bengio, *Practical Recommendations for Gradient-Based Training of Deep Architectures*, 2012. [PDF](#)
- Goodfellow, Bengio, Courville, *Deep Learning*, 2016 [Website](#)

# Outline

- Introduction
- Supervised Learning
- Unsupervised Learning
- Applications to Computer Vision
- Conclusions
- Machine Learning for Drones

# The drone market is valued \$24 billions today

Inspection



Agriculture



Transport



Search and Rescue



(c) image is copyright protected - www.asctec.de

Source: Swiss Drone Industry Report 2021, p. 22:

[https://drive.google.com/file/d/1ljesolDoUu1-IVX14nqJRCT-wpEQB22 /view](https://drive.google.com/file/d/1ljesolDoUu1-IVX14nqJRCT-wpEQB22/view)

# How are current commercial drones controlled?

- **By a human pilot**
  - requires **line of sight** or **video link**
  - requires **a lot of training**
- **By an autopilot:** autonomous navigation
  - **GPS:** doesn't work in GPS denied or degraded environments
  - **Lidar** (e.g., Exyn): expensive, heavy, power hungry
  - **Cameras** (e.g., Parrot, DJI, Skydio): cheap, lightweight, passive (i.e., low power)



# Last 10-years Progress on Autonomous Vision-based Flight

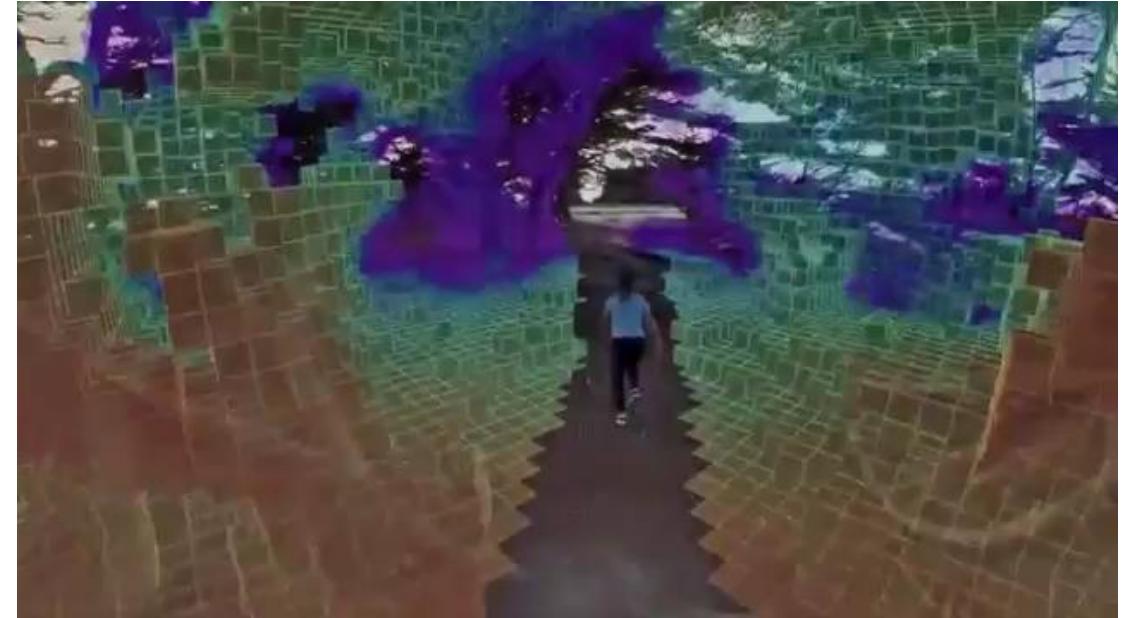


**2010**

**EU SFLY Project (2009-2012)**

[[Bloesch, ICRA 2010](#)]

**1<sup>st</sup> onboard goal-oriented  
vision-based flight**  
(previous research focused  
on reactive navigation)

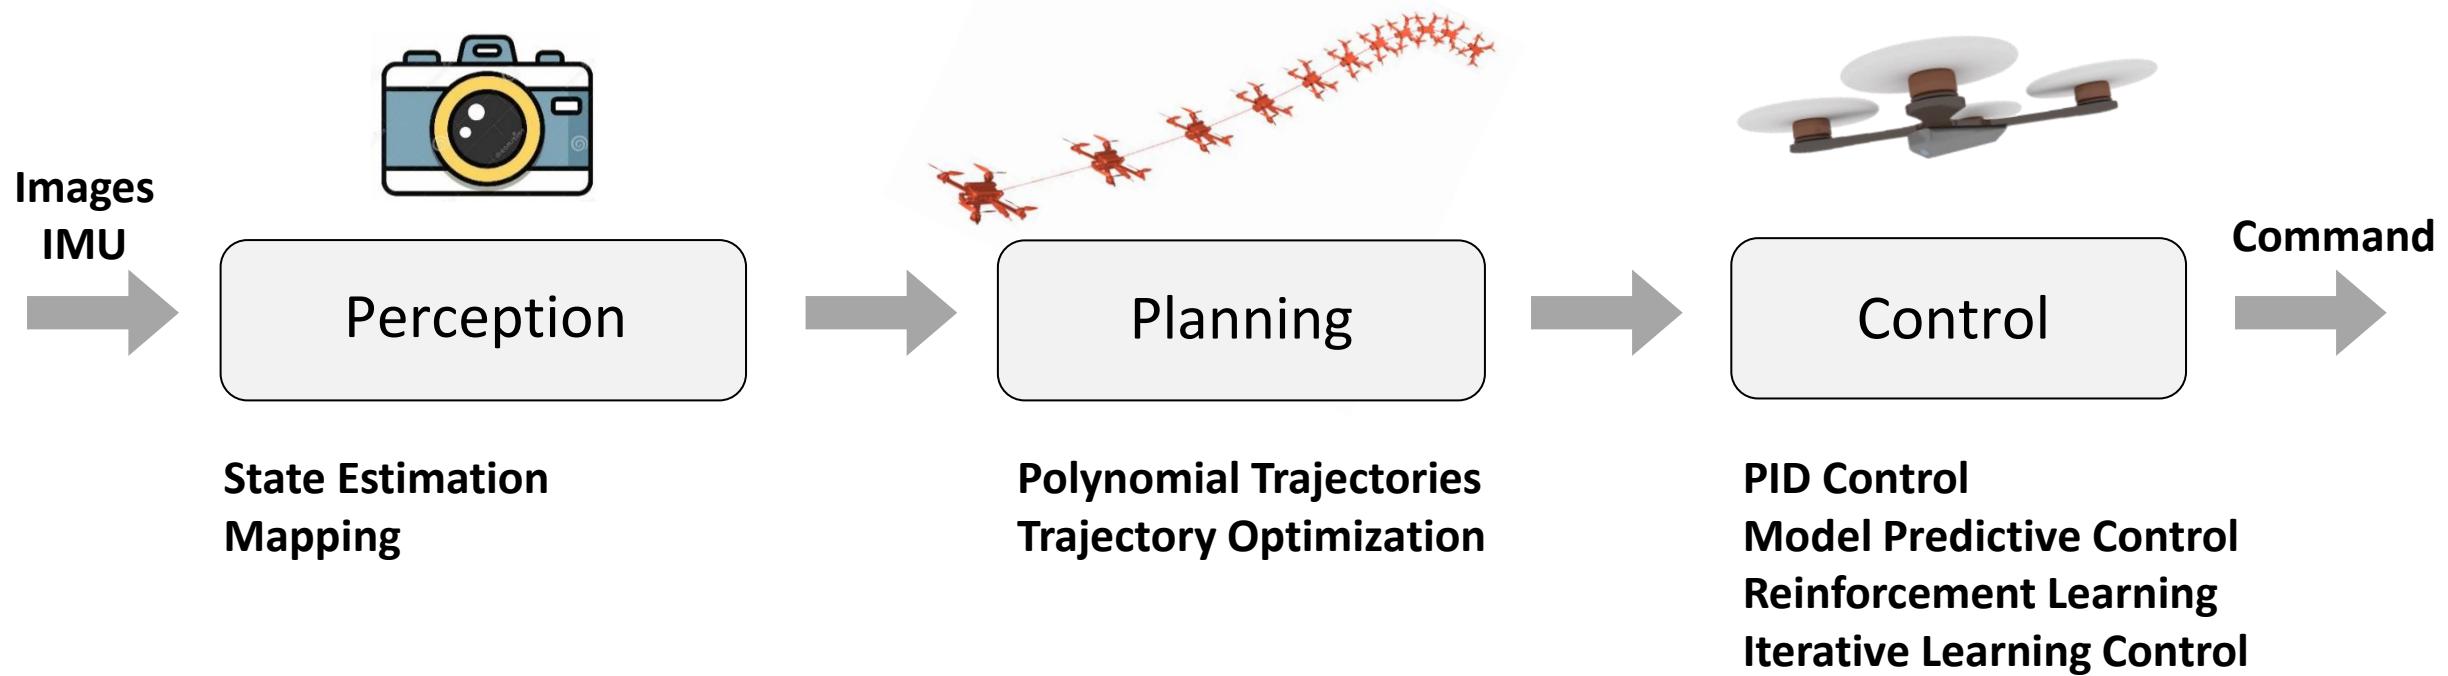


**2020**

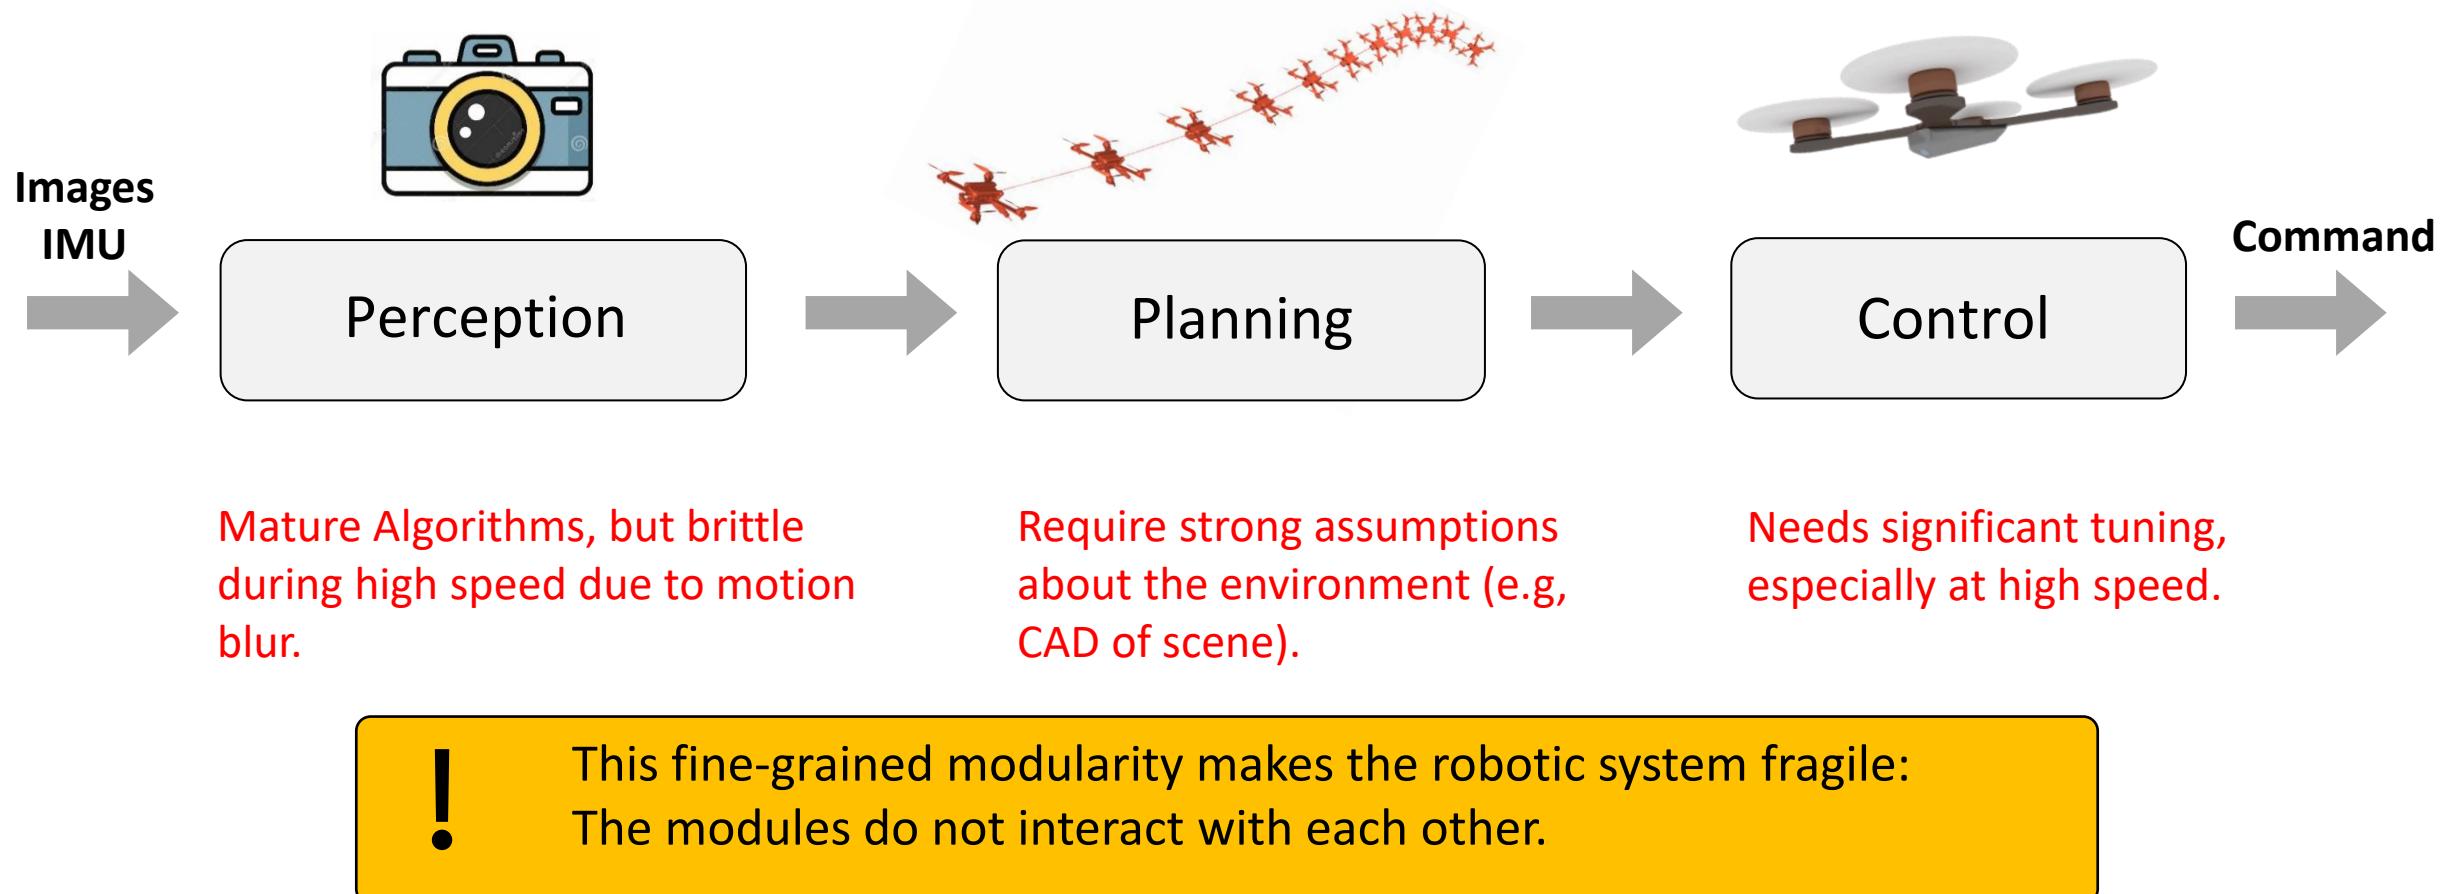
- **Skydio (2018-2020),**
- **DJI (2018-2020),**
- **NASA Mars Helicopter (2020)**

**1<sup>st</sup> products in the market  
or sent to another planet ☺**

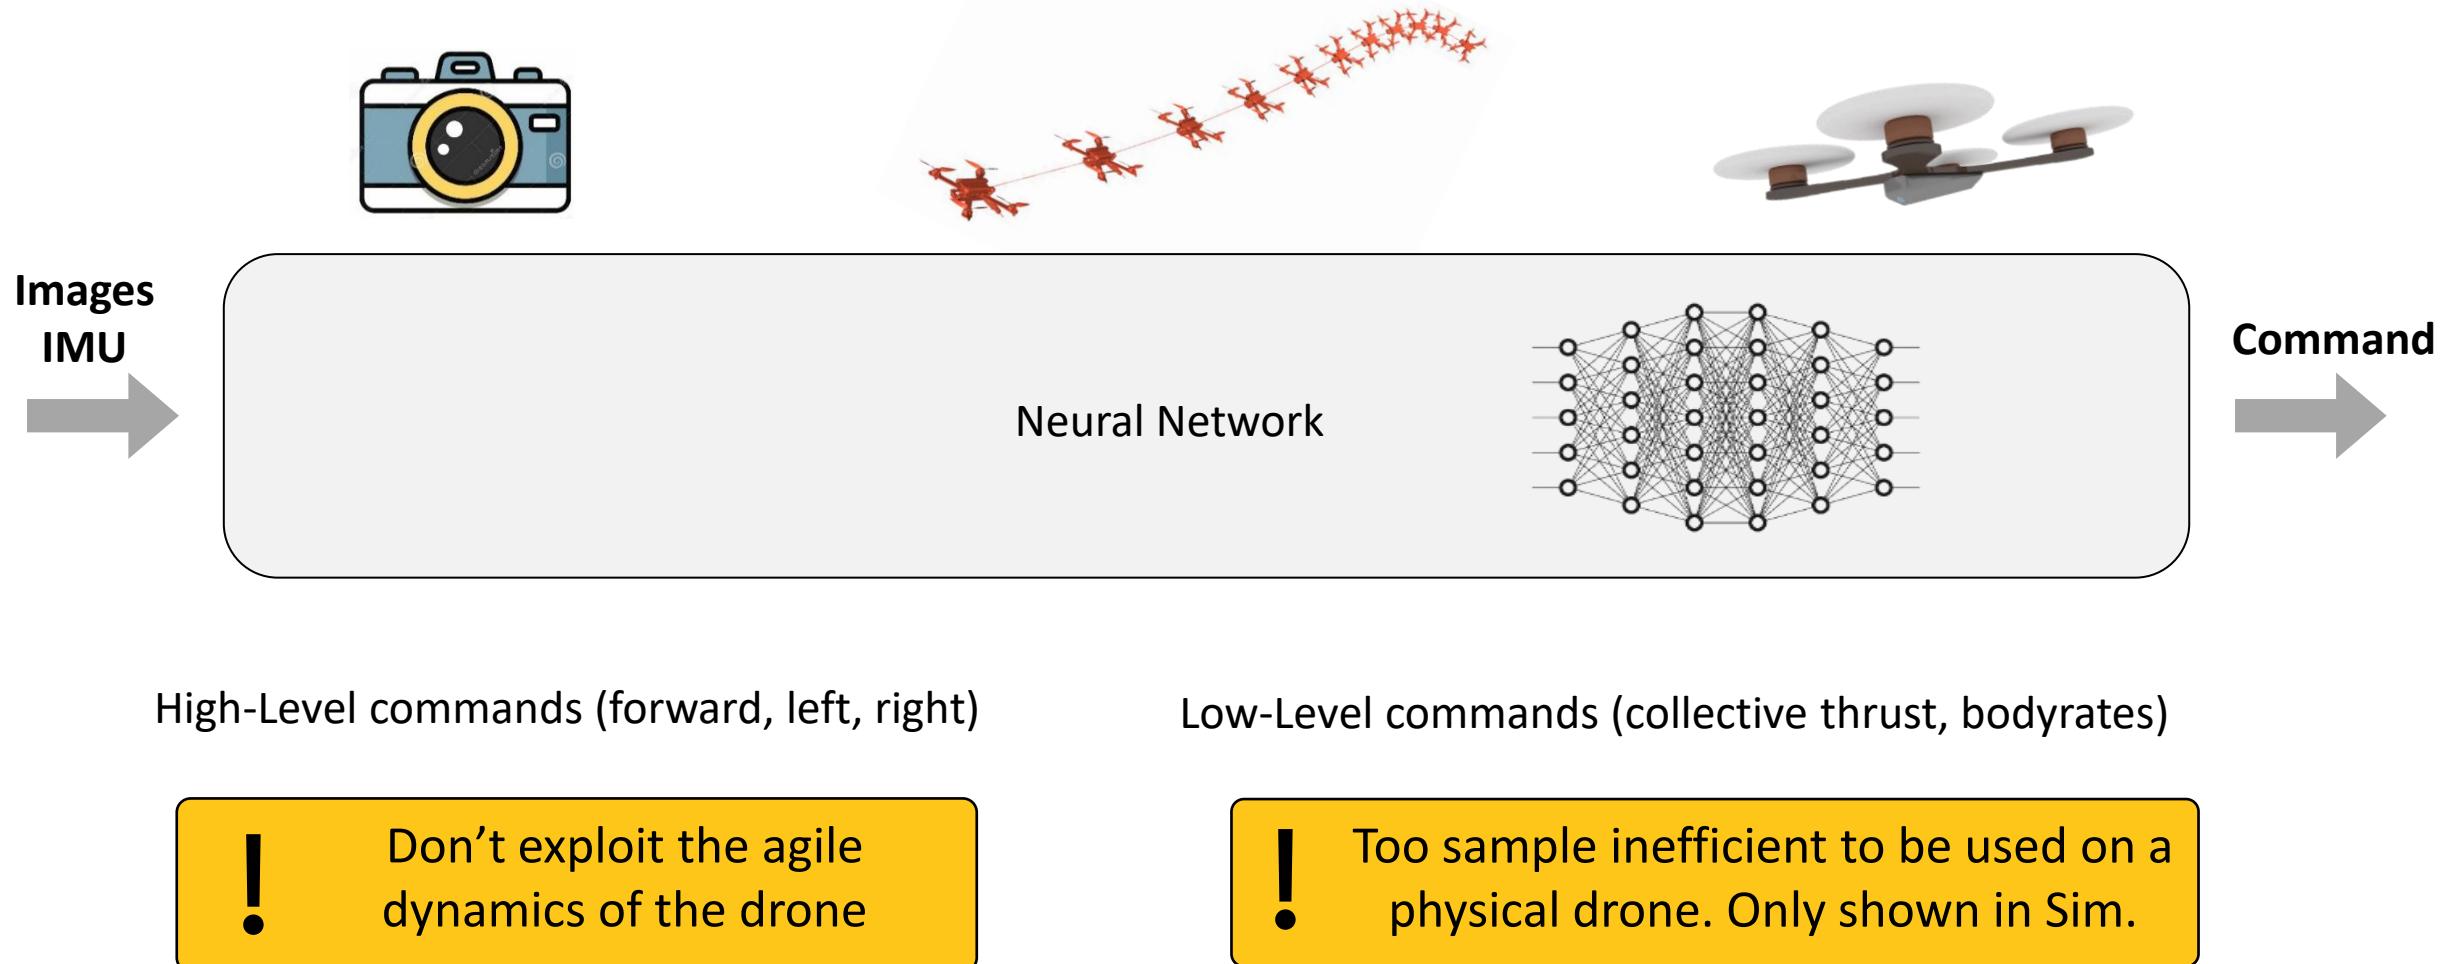
# Related Work: The Traditional Approach



# Related Work: The Problem



# Related Work: End-to-End Learning



# Our Research



Augment the traditional robotic cycle  
with learning-based methods.

## Hypothesis:

Neural Networks can distill the knowledge of mature robotics algorithms  
into computationally efficient and robust sensorimotor policies.

# Projects

- Learning High-Speed Flight in the Wild



- Autonomous Drone Racing



# Learning High-Speed Flight in the Wild

What does it take to achieve similar **spatial awareness** to a human **with comparable sensing (and computing)** in the context of **high-speed flight**?

Assumptions:

Human pilots fly under similar assumptions!

- No external sensing or computing.
- Test environment not seen in advance.
- Possibly dynamic environment.

Available Information:

- Visual Feedback (multiple cameras).
- Inertial Feedback.
- An intention (e.g. fly straight).



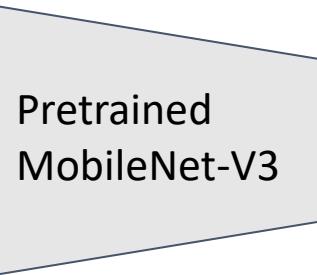
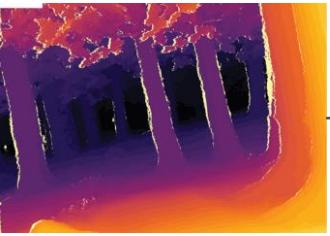


Loquercio, Kaufmann, Ranftl, Mueller, Koltun, Scaramuzza: Learning, High-Speed Flight in the Wild,  
Science Robotics, 2021. [PDF](#), [Video](#), [Code](#)

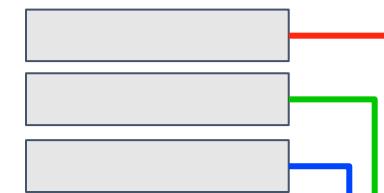
# Multiple-Hypothesis Action Prediction

We predict collision-free **receding-horizon trajectories** using a **neural network** with access to visual and inertial observations, as well as a reference velocity.

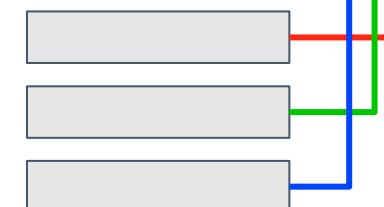
Stereo Depth



$M = 3$

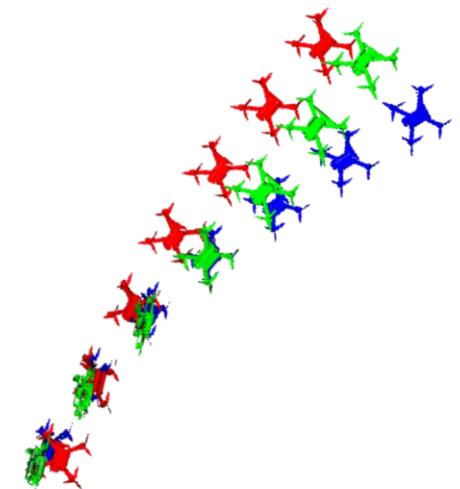


$M \times 32$

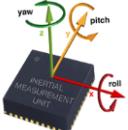


$M \times 32$

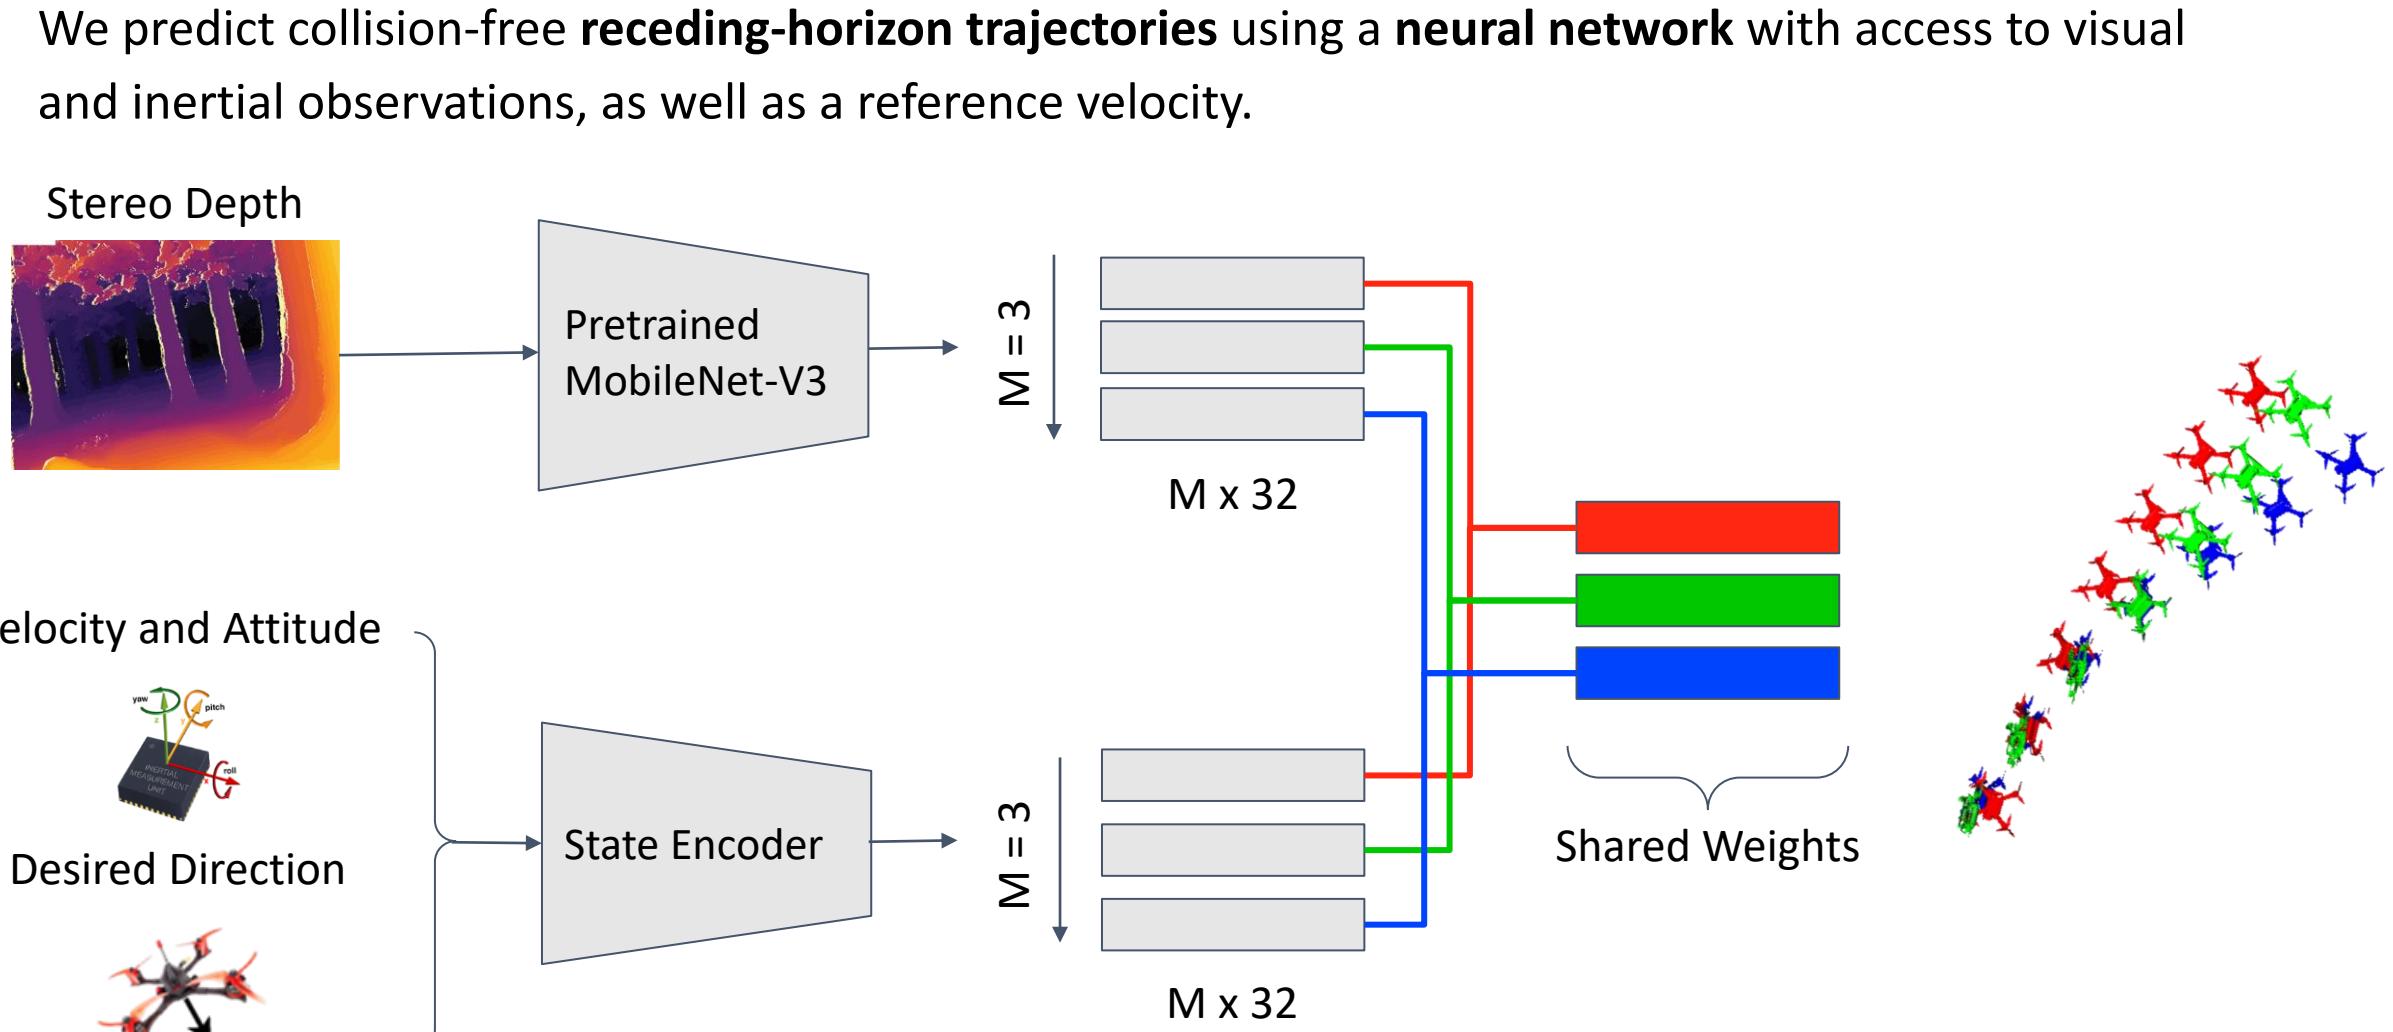
Shared Weights



Velocity and Attitude



Desired Direction



# Training Procedure

We follow the **privileged learning paradigm** to train the network **purely in simulation\***.

1. **Design an expert planner with access to full knowledge of the environment.**

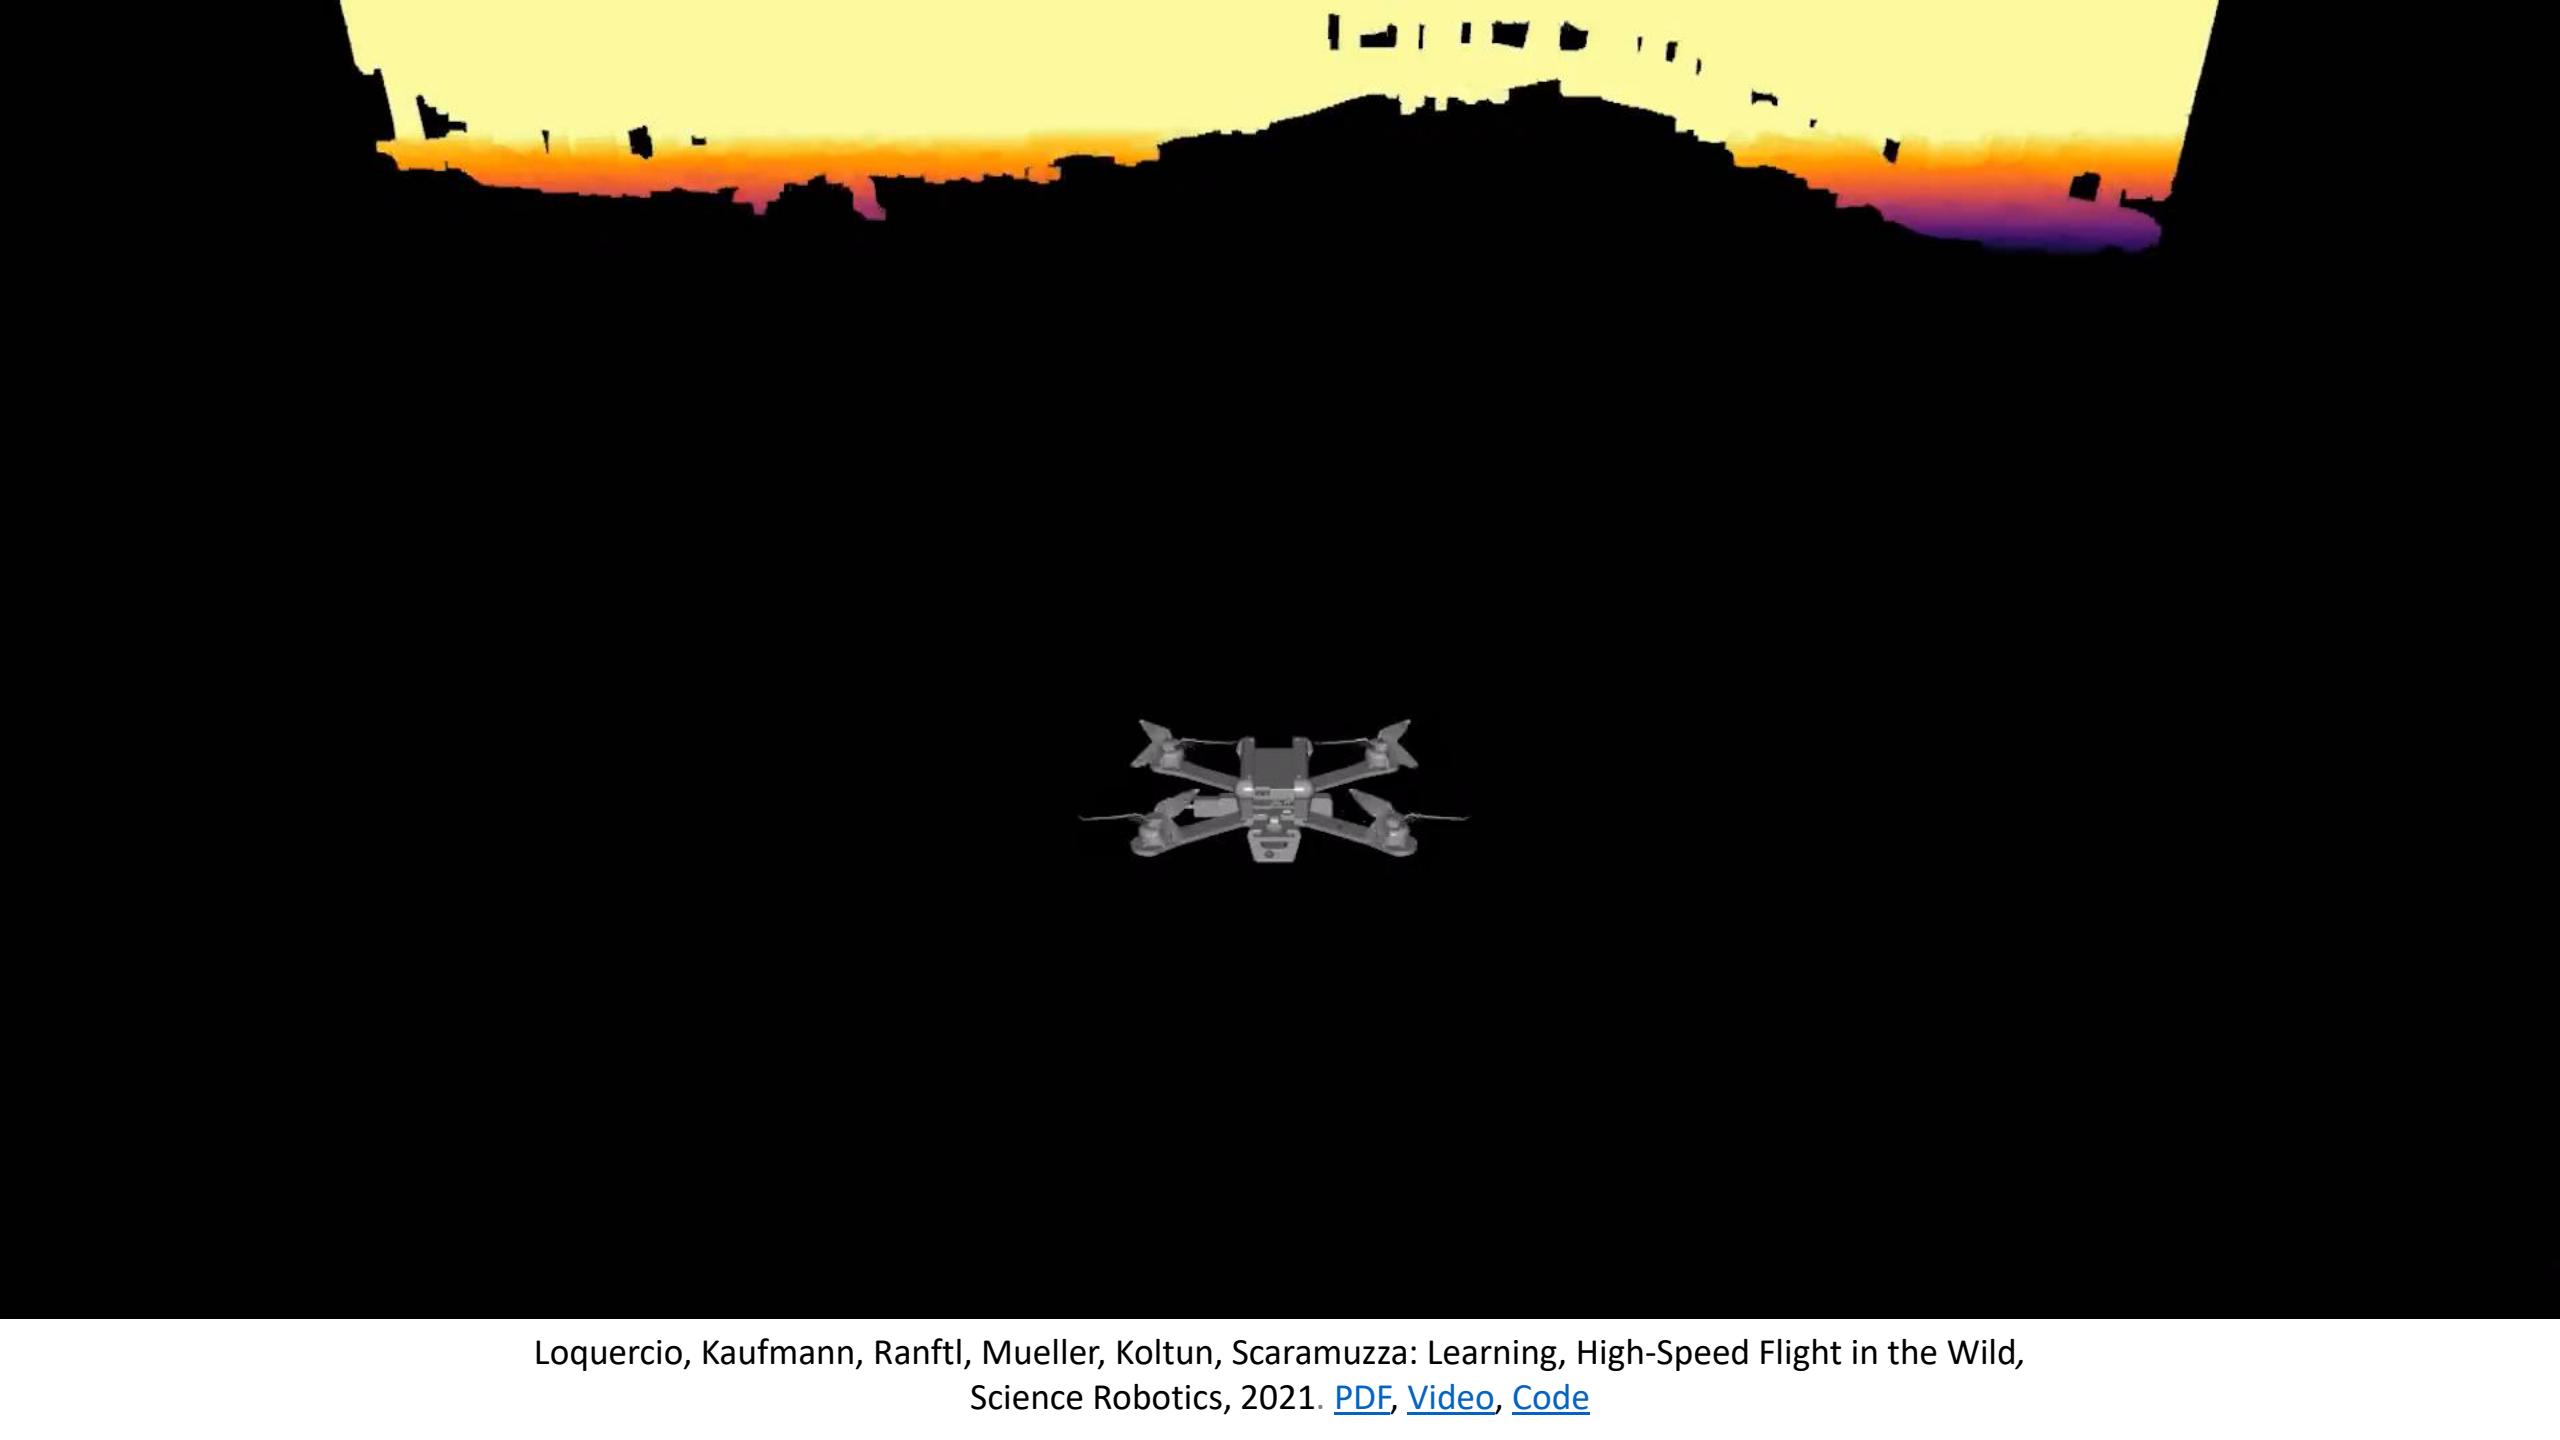
This expert uses a fine-grained point-cloud of the scene to find collision-free trajectories with sampling.

2. **Distill the knowledge of the expert into a deep neural network.**

Basically do imitation learning from a set of expert demonstrations.

This simple idea hides quite some challenges!

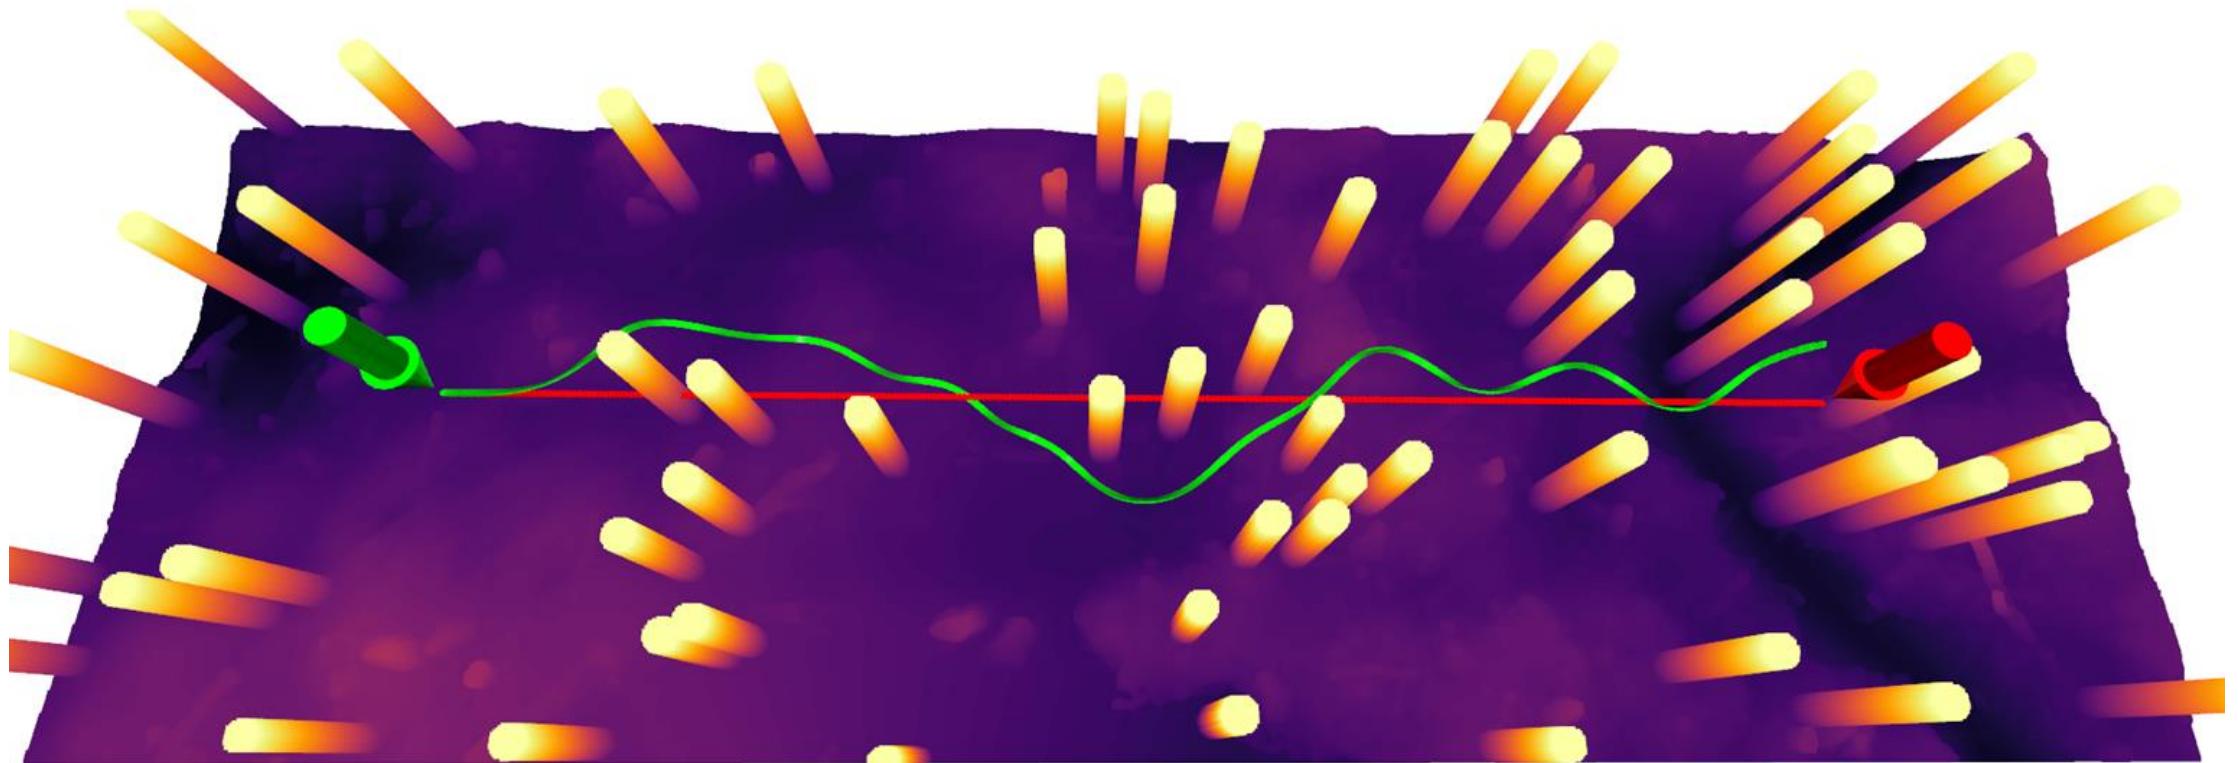
\* Impossible to collect a dataset of real-world demonstrations since it is not possible (or very expensive) to have a perfect map of the environment.



Loquercio, Kaufmann, Ranftl, Mueller, Koltun, Scaramuzza: Learning, High-Speed Flight in the Wild,  
Science Robotics, 2021. [PDF](#), [Video](#), [Code](#)

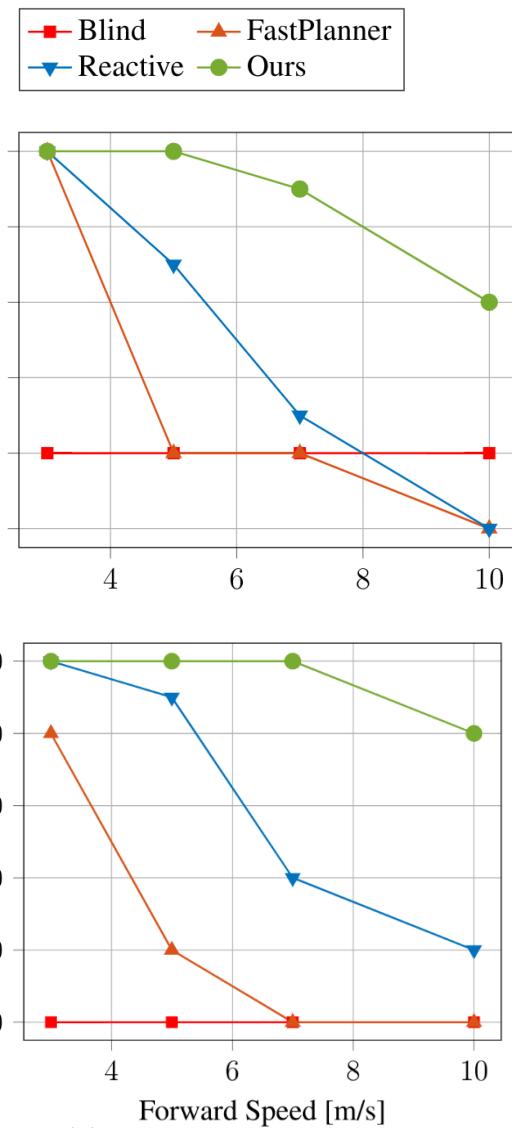
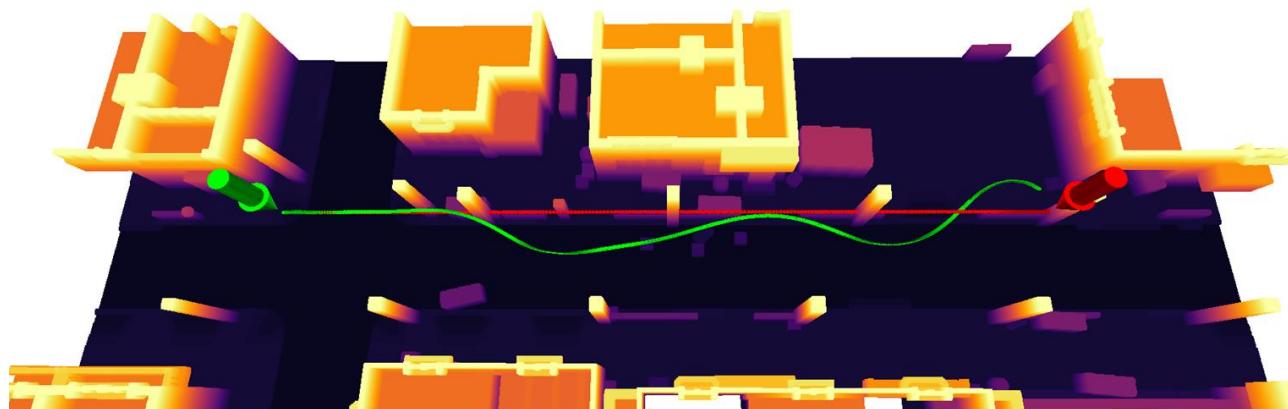
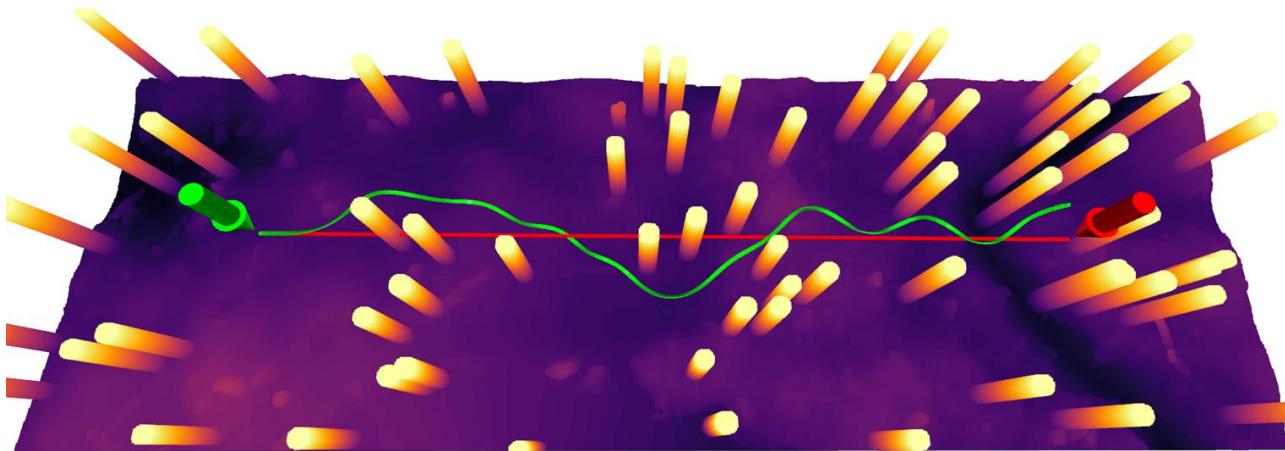
# Controlled Experiments

Evaluate on the task of reaching a goal with no prior knowledge about the scene.



# Controlled Experiments

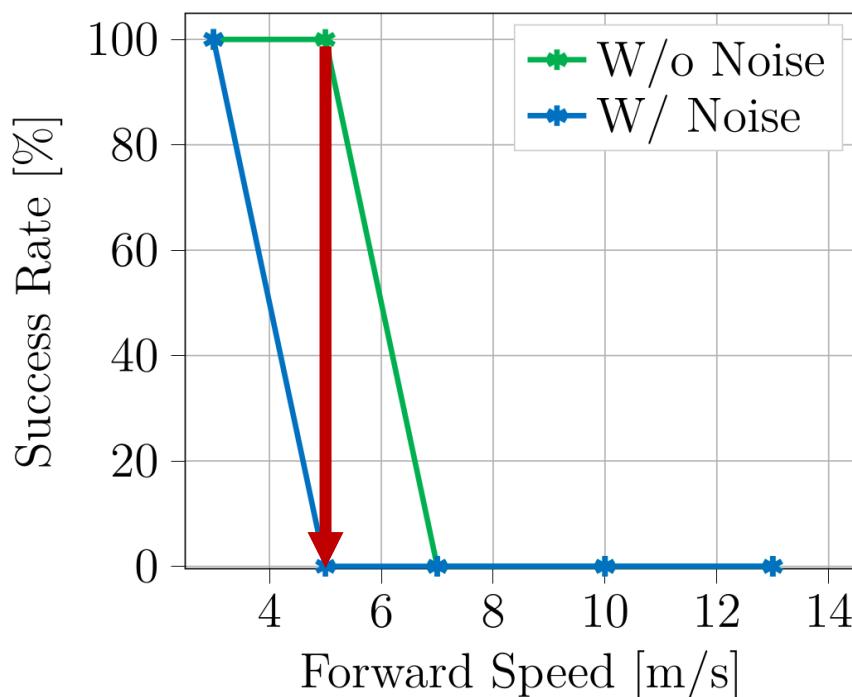
Comparison with state-of-the-art navigation methods.



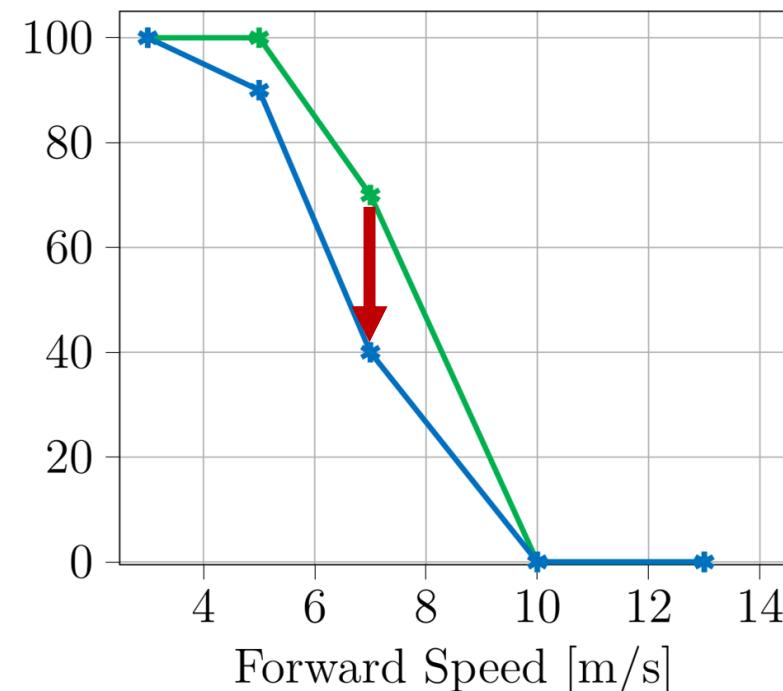
# Controlled Experiments

Robustness to sensor noise:

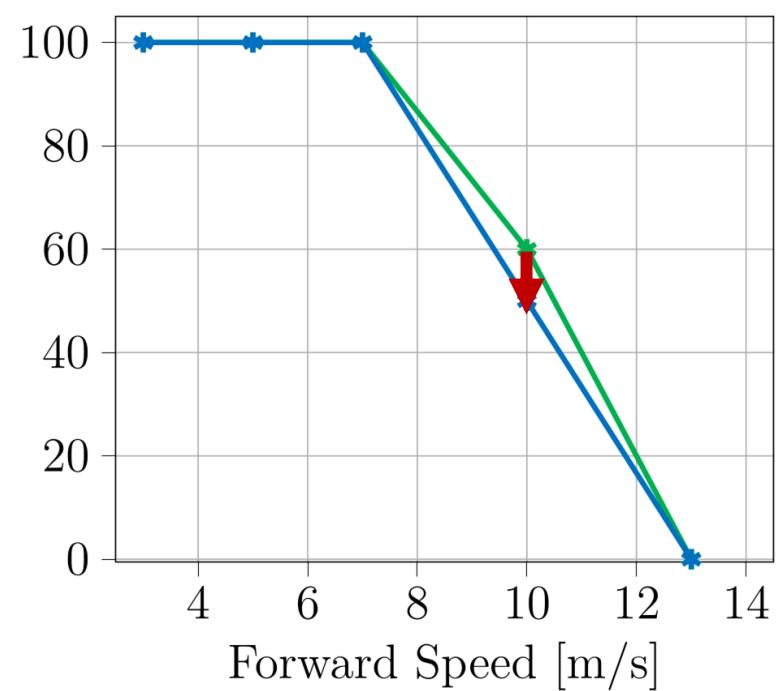
Zhou et al., TRO-2020

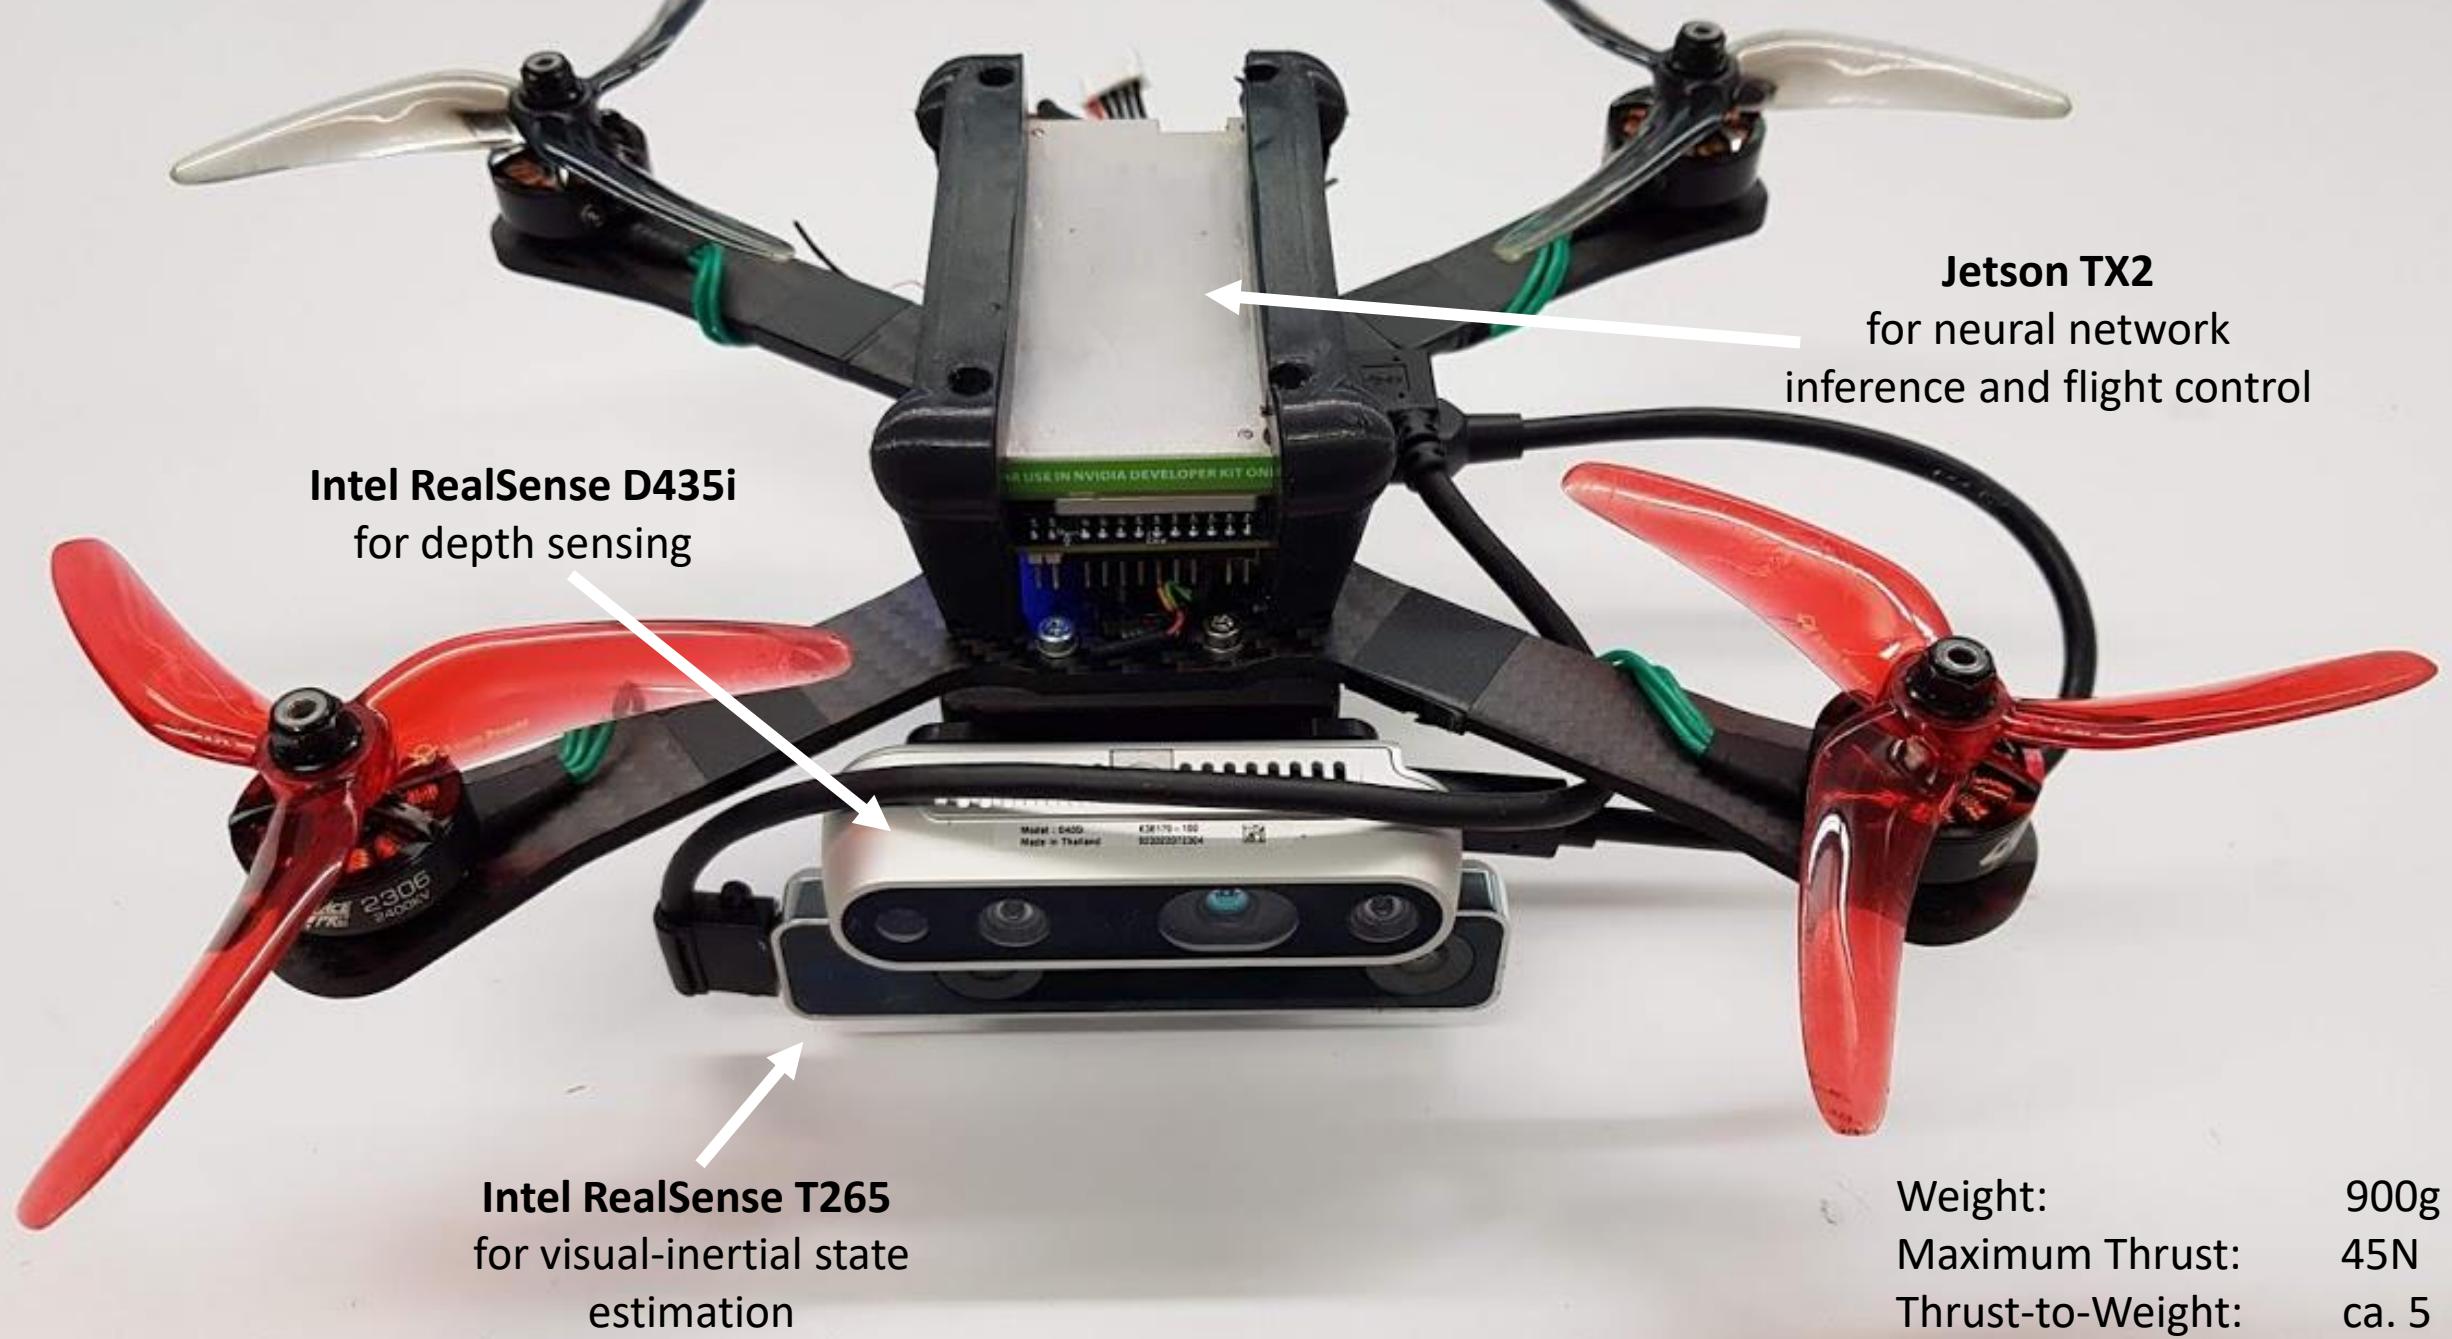


Florence et al., AFR-2020



Kaufmann et al. (Ours)







# Autonomous Drone Racing



**WARNING!** This drone is NOT autonomous; it is operated by a human pilot.

**Human pilots take years** to acquire the skills shown in this video.

# Challenges of Autonomous Drone Racing

- **Real-time coupling of perception and action**
- Coping with **inaccurate models of sensors, actuators, environment**
- Coping with **dynamically changing** environments
- Coping with **unreliable perception and state estimation:**
  - low texture
  - HDR scenes
  - motion blur



# Autonomous Racing in Motion Capture



- [1] Foehn et al., *Time-Optimal Planning for Quadrotor Waypoint Flight*, *Science Robotics*, 2021. [PDF](#). [Video](#). [Code](#). Featured on [Forbes](#) magazine.
- [2] Foehn et al., *AlphaPilot: Autonomous Drone Racing*, *RSS 2020*, Best Systems Paper Award. [PDF](#) [Video](#). 2<sup>nd</sup> place at AlphaPilot Challenge
- [3] Kaufmann et al., *Beauty and the Beast: Optimal Methods Meet Learning for Drone Racing*, *ICRA'19*. [PDF](#). [Video](#). 1<sup>st</sup> place at IROS'18 Drone race. [Video](#).
- [4] Loquercio, et al., *Deep Drone Racing*, *IEEE Transactions on Robotics* 2020. Best Paper Award finalist. [PDF](#). [Video](#)
- [5] Song et al, *Autonomous Drone Racing with Deep Reinforcement Learning*, *IROS'21*. [PDF](#). [Video](#)
- [6] Simulator used: Song et al., *Flightmare: A Flexible Quadrotor Simulator*, *CORL'20*, [PDF](#) [Video](#) [Website](#)

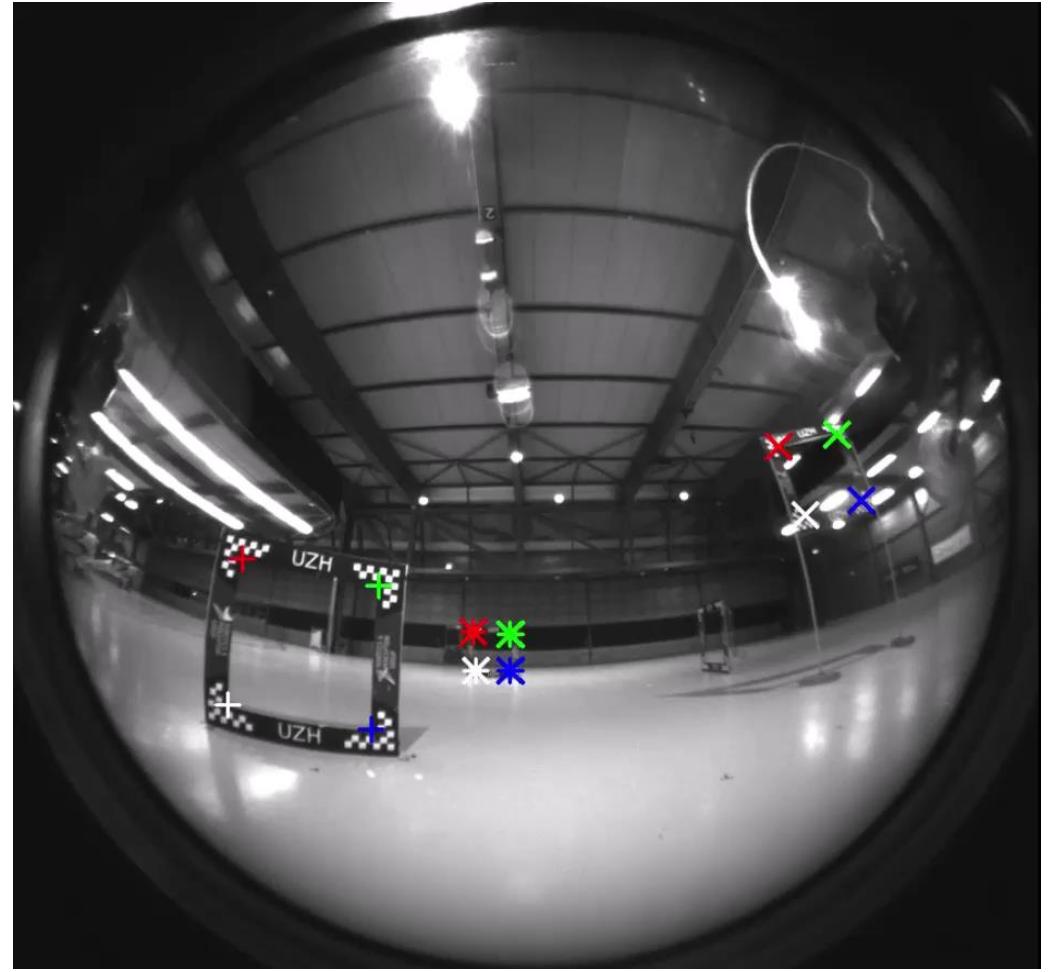
# Autonomous Racing with Onboard Sensing



- [1] Foehn et al., *Time-Optimal Planning for Quadrotor Waypoint Flight*, *Science Robotics*, 2021. [PDF](#). [Video](#). [Code](#). Featured on [Forbes](#) magazine.
- [2] Foehn et al., *AlphaPilot: Autonomous Drone Racing*, *RSS 2020*, Best Systems Paper Award. [PDF](#) [Video](#). 2<sup>nd</sup> place at AlphaPilot Challenge
- [3] Kaufmann et al., *Beauty and the Beast: Optimal Methods Meet Learning for Drone Racing*, *ICRA'19*. [PDF](#). [Video](#). 1<sup>st</sup> place at IROS'18 Drone race. [Video](#).
- [4] Loquercio, et al., *Deep Drone Racing*, *IEEE Transactions on Robotics* 2020. Best Paper Award finalist. [PDF](#). [Video](#)
- [5] Song et al, *Autonomous Drone Racing with Deep Reinforcement Learning*, *IROS'21*. [PDF](#). [Video](#)
- [6] Simulator used: Song et al., *Flightmare: A Flexible Quadrotor Simulator*, *CORL'20*, [PDF](#) [Video](#) [Website](#)

# Autonomous Racing with Onboard Sensing

- VIO estimates are combined with learning-based gate detections to compensate for drift.
- VIO, neural network inference are running onboard



[1] Foehn et al., *Time-Optimal Planning for Quadrotor Waypoint Flight*, *Science Robotics*, 2021. [PDF](#). [Video](#). [Code](#). Featured on [Forbes](#) magazine.

[2] Foehn et al., *AlphaPilot: Autonomous Drone Racing*, *RSS 2020*, Best Systems Paper Award. [PDF](#) [Video](#). 2<sup>nd</sup> place at AlphaPilot Challenge

[3] Kaufmann et al., *Beauty and the Beast: Optimal Methods Meet Learning for Drone Racing*, *ICRA'19*. [PDF](#). [Video](#). 1<sup>st</sup> place at IROS'18 Drone race. [Video](#).

[4] Loquercio, et al., *Deep Drone Racing*, *IEEE Transactions on Robotics* 2020. Best Paper Award finalist. [PDF](#). [Video](#)

[5] Song et al., *Autonomous Drone Racing with Deep Reinforcement Learning*, *IROS'21*. [PDF](#). [Video](#)

[6] Simulator used: Song et al., *Flightmare: A Flexible Quadrotor Simulator*, *CORL'20*, [PDF](#) [Video](#) [Website](#)

# Conclusions and Takeaways

- Autonomous **vision-based agile flight** as a new research topic (at least 10 years to solve it)
  - **Pushes the limit of existing algorithms** in extreme situations
  - Raises **fundamental problems** for robotics **research**



# Check out our student projects!

- Visit our webpage: [http://rpg.ifi.uzh.ch/student\\_projects.php](http://rpg.ifi.uzh.ch/student_projects.php)

## Reinforcement Learning for Offboard Control of a Racing Drone - [Available](#)



**Description:** Autonomous drone racing using offboard control is very challenging because a limited amount of sensor data are available to the offboard computer. Offboard control however, allows using extremely lightweight drones that can reach much higher performances than heavier drones using onboard control. This project has two goals: First, implement and test a communication interface between a ROS-based flight stack and a C/Python-based codebase for low-latency high-bandwidth communication between the drone and an offboard computer. Second,

use reinforcement learning to develop a policy that can successfully fly an offboard-controlled racing drone in the real world. Requirements: Strong background in robotics and machine learning is required. ROS, C++, and Python skills. Experience with quadrotor hardware and drone flight is a plus.

## Perception Aware Minimum-time Planning in Cluttered Environments - [Available](#)



**Description:** Autonomous drone racing requires planning trajectories that pass through a sequence of gates as fast as possible to beat other competitors. However, planning high-speed trajectories in obstacle-cluttered environments for dynamic systems like drones is a challenging problem. Moreover, the planning algorithms should support autonomous vision-based flight by creating trajectories with respect to the onboard camera used for both state estimation and gate detection. The goal of this project is to develop planning algorithms for drones that consider known cluttered environments and plan perception aware minimum-time trajectories. Applications should have strong experience in C++ and experience with ROS.

## Benchmarking Algorithms for Autonomous Drone Racing - [Available](#)



**Description:** Drone racing is an increasingly popular sport and many world-class pilots use simulators for practicing their skills. This project aims to deploy control algorithms for autonomous drone racing in the Liftoff Drone Racing simulator (<https://www.liftoff-game.com/>), collect data from the autonomous and human-piloted flight, and perform a benchmark comparison between both. The student will implement a control interface between the Liftoff simulator and an existing ROS-

based flight stack, collect and analyze flight performance data. Requirements: Strong programming experience in ROS, C++, and Python. Experience in Unity3D is a plus.

## Deep Learning for Model Predictive Contouring Control - [Available](#)



**Description:** Model Predictive Contouring Control (MPCC) has shown to achieve very good results in the task of time-optimal multi-waypoint flight. MPCC methods have the freedom to select the optimal states of the system at runtime, dropping the need for a computationally expensive reference trajectory. Our recent work shows MPCC can achieve better lap times than state-of-the-art planning+tracking approaches, and that the method can be run in real-time.

**Goal:** One of the extra benefits of the MPCC approach is that there are only two relevant parameters to be tuned in the cost function: contour weight and progress weight. In this project, we aim to exploit the low dimensionality of this tuning parameter space and apply learning techniques to find a mapping from a high-level task (track waypoints in a certain order in minimum time, for example) to MPCC tuning parameters.